

CUBIQO — Master Strategic Report (v2)

Prepared by: MO (CTO / AI Co-Founder)
Date: 2026-02-22 (updated)
Confidential — Founder Eyes Only

*This document consolidates all analyses conducted on 2026-02-21/22.
Every finding is grounded in direct inspection of the actual codebase:
157 API routes · 44 database migrations · 80 UI components.
4,586 lines · Sections 1-8 · Appendices A-E · 20 Techno-Functional Requirements*


EXECUTIVE SUMMARY — Cubiqo Product State (February 2026)

Prepared by: MO — CTO / AI Co-Founder
Based on: Direct code inspection of branch `copilot/investigate-features-and-ui-components` + full analysis corpus (8 documents, 3,100+ lines)

What Cubiqo Is

Cubiqo is an AI-native operating system for solopreneurs and knowledge workers. It routes voice and text input through a policy-aware LLM backend (TEAL/RED/YELLOW zones), orchestrates a suite of integrated tools (Social Army, Journal, Job Hunt, Emergent code editor, RGY peer matching, BYO API keys), and presents itself through a 3D animated cuboid UI that reflects AI state in real time.

Overall Product State at a Glance



3 Absolute Launch Blockers (Cannot Ship Without These)

#	Blocker	Why Critical	Est. Effort
1	Stripe billing + CubiKey payment	CubiKey is placeholder UI — no revenue collection possible	3 weeks
2	Terms of Service + Privacy Policy pages	GDPR/CCPA requires these before any user data is collected	3 days
3	Age gate for RED zone	Legal liability — minors can access age-gated/ explicit content	1 week

Current verdict: Cubiqo is NOT legally launchable today. These 3 items must ship first.

What Is Genuinely Production-Ready (Ship Now)

Area	Confirmation
Policy Router (LLM backend, zone detection, failover, safety)	<code>src/lib/ai/policy-router.ts</code> — solid, 3-backend failover
Color system (TEAL/RED/YELLOW, voice tones, lock commands)	<code>src/config/colors.ts</code> + <code>commands.ts</code>
STT + TTS pipeline (Whisper + ElevenLabs + voice modulation)	<code>/api/stt/</code> , <code>/api/tts/</code> , <code>voice-modulation.ts</code>
Auth (magic link, OAuth, WebAuthn/Passkeys)	4 auth modes fully wired
BYO Mode (cloud API keys, AES-256 encryption, routing)	<code>byo-manager.ts</code> , <code>/api/byo/</code>
Browser automation (Puppeteer, consent gate)	<code>BrowserService.ts</code> , <code>/api/browser/</code>
RGY capsule matching (4-signal: colour→intent→keyword→vector)	<code>capsule-manager.ts</code> , <code>discovery-service.ts</code>
Side Panel keywords (per-color RGY lists, session persistence)	<code>KeywordPanel.tsx</code>
3D cuboid (wireframe + solid, 5 Special Moves, state-reactive)	<code>AICuboidGLB.tsx</code> , <code>SilverWireLandingCube.tsx</code>
Conscious memory (extraction, storage, journey tracking)	<code>/api/memory/</code> , <code>/api/extract-memories/</code>
Emergent code editor (9 API modules, Monaco, sandboxed preview)	<code>/api/emergent/</code> routes
Social Army (10 platforms, persona accounts, GFXToolz)	<code>social-army-service.ts</code> , <code>platforms.json</code>
Journal / Rozana (BigBoss prompts, 24h gate, color classification)	<code>/api/journal/</code> , <code>journal-service.ts</code>

Pending Work Summary (All Priorities)

Priority	Count	Items
● P0 — Launch Blockers	5	Stripe, ToS/PP, Age Gate, Zero-Retention clarification, Smart-home removal or stub
● P1 — 30-Day Post-Launch	10	CAP Orchestrator, Audio Cues, Vocspad, 3 Special Moves, CQ↔CQ UI, Spending Caps persistence, Analytics funnel, Social Army review gate, Onboarding branching
● P2 — 90-Day Enhancement	10	Calendar API, Wallet/Stripe crypto, Geo-fence UI, SettingsCube spoken confirm, Fabric cube material, CQ Score, Food/taxi live API, Referral programme, Emergent collaborative editing, Structured event logs

Key Numbers From Code (Ground Truth)

Metric	Value	Source
Total API routes	157	<code>find /api -name route.ts</code>
Disabled/stub routes	17	<code>// TODO , return NextResponse.json({message: 'coming soon'})</code>
TODO / stub markers	101+	<code>grep -r "TODO\ FIXME\ STUB"</code>
Analytics events tracked	3	<code>trackEvent()</code> call sites
Stripe references	0	No <code>stripe</code> in <code>package.json</code>
Social platforms configured	10	<code>platforms.json</code>
DB migrations applied	18	<code>supabase/migrations/</code>
Patent opportunities (≥50% approval)	4	See <code>PATENT_OPPORTUNITIES.md</code>
Spec compliance (Flagship doc)	63%	See Section 11 below

TECHNO-FUNCTIONAL REQUIREMENTS — All Incomplete & Pending Features

Format for each TFR:

- **Current State** — what exists in code right now
- **Requirement** — what the feature must do when complete
- **Technical Approach** — specific implementation path
- **Changes Needed** — DB schema / API endpoints / UI components
- **Acceptance Criteria** — how to know it's done

- **Effort Estimate** — realistic solo developer estimate
- **Owner** — which agent/role should implement it

TFR-001 — Stripe Billing + CubiKey Subscription Payment

Priority: 🟡 P0 — Launch Blocker

Current State:

- `src/components/CubiKey.tsx` exists but is placeholder UI only
- `src/lib/cubikey/cubikey-manager.ts` references Stripe but no `stripe` package in `package.json`
- No `/api/billing/` routes exist
- No `subscription_plans` or `user_subscriptions` DB table

Requirement:

Users must be able to purchase CubiKey Pro (\$29/mo) and CubiKey Max (\$79/mo) via Stripe checkout. Subscription status must gate features in real time.

Technical Approach:

1. `npm install stripe @stripe/stripe-js`
2. Create Stripe products + price IDs in Stripe dashboard
3. Build `/api/billing/checkout` → creates Stripe Checkout Session
4. Build `/api/billing/webhook` → handles `subscription.created/updated/deleted`
5. Build `/api/billing/portal` → Stripe Customer Portal for self-serve management
6. Add `user_subscriptions` table to Supabase
7. Add feature gate check to `policy-router.ts` (plan tier → model access)
8. Replace `CubiKey.tsx` placeholder with real checkout button

DB Changes:

```
CREATE TABLE user_subscriptions (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  user_id UUID REFERENCES auth.users(id) ON DELETE CASCADE,  
  stripe_customer_id TEXT UNIQUE,  
  stripe_subscription_id TEXT UNIQUE,  
  plan_id TEXT NOT NULL,           -- 'free' | 'pro' | 'max'  
  status TEXT NOT NULL,           -- 'active' | 'past_due' | 'canceled'  
  current_period_end TIMESTAMPTZ,  
  created_at TIMESTAMPTZ DEFAULT now(),  
  updated_at TIMESTAMPTZ DEFAULT now()  
);  
CREATE INDEX ON user_subscriptions(user_id);  
CREATE INDEX ON user_subscriptions(stripe_customer_id);
```

API Endpoints:

- POST `/api/billing/checkout` — create Stripe Checkout session, return URL
- POST `/api/billing/webhook` — Stripe webhook handler (must be raw body)
- GET `/api/billing/portal` — create Stripe Customer Portal session
- GET `/api/billing/status` — return current plan for authenticated user

UI Components:

- `CubiKeyCheckoutButton.tsx` — Stripe Checkout trigger

- `BillingPortalLink.tsx` — link to manage subscription
- `PlanGateBadge.tsx` — upgrade prompt when feature is gated

Acceptance Criteria:

- [] Test card `4242 4242 4242 4242` completes checkout in staging
- [] Webhook updates `user_subscriptions.status` within 5 seconds
- [] Free user cannot access `plan_id='pro'` features
- [] Pro user upgrade/downgrade visible in Stripe dashboard

Effort: 3 weeks (1 week Stripe setup + webhook; 1 week feature gates; 1 week UI)

Owner: Blossom (backend) + Bubbles (UI)

TFR-002 — Terms of Service + Privacy Policy Pages

Priority: ● P0 — Launch Blocker

Current State:

- No `/terms` or `/privacy` route exists anywhere in the app
- No ToS or Privacy Policy text in the repository
- `magic-link.ts` email template says "Zero-Retention. Private." — contradicts memory storage

Requirement:

Legal pages must be live before any user data is collected. Must cover: data storage (memory system), AI processing (third-party LLMs), age restrictions (RED zone), subscription terms, and GDPR/CCPA rights.

Technical Approach:

1. Create `src/app/terms/page.tsx` — static MDX page, dated
2. Create `src/app/privacy/page.tsx` — static MDX page, GDPR-compliant
3. Add footer links to both pages across all layouts
4. Add checkbox "I agree to Terms + Privacy Policy" to signup form
5. Store consent timestamp to user profile (GDPR requirement)
6. Fix `magic-link.ts` email to remove "Zero-Retention" (false claim)

DB Changes:

```
ALTER TABLE user_profiles ADD COLUMN IF NOT EXISTS
  terms_accepted_at TIMESTAMPTZ,
  terms_version TEXT DEFAULT '1.0';
```

10 Required ToS Clauses (minimum):

1. Subscription terms, refund policy, cancellation
2. AI-generated content disclaimer (not professional advice)
3. Age restriction (18+ for RED zone, 13+ general)
4. Data storage acknowledgement (memory system, Supabase)
5. Third-party LLM processing (OpenAI, Anthropic, ElevenLabs)
6. BYO API keys: user owns and is responsible for their keys
7. Social Army: user is responsible for content posted via their accounts
8. Acceptable use policy (no abuse, spam, illegal content)
9. GDPR/CCPA data subject rights (access, deletion, portability)
10. Limitation of liability + governing law (Ontario, Canada)

Acceptance Criteria:

- [] `/terms` and `/privacy` render as static pages with correct content
- [] Signup flow cannot complete without ToS checkbox checked
- [] `terms_accepted_at` is populated for new users
- [] Footer links present on landing, dashboard, and onboarding pages

Effort: 3 days (legal text from Termly.io template; 1 day engineering)

Owner: MO (content) + Bubbles (UI wiring)

TFR-003 — Age Gate for RED Zone

Priority: ● P0 — Launch Blocker

Current State:

- `src/config/colors.ts` defines RED zone as `emotion: 'age-gated, critical, goal-oriented'`
- `policy-router.ts` routes RED zone to uncensored models (`MIXTRAL_8X22B` , `LLAMA_UNCENSORED`)
- Zero age verification code exists anywhere in the codebase
- Any user can activate RED zone through the color selector

Requirement:

Users must confirm they are 18+ before accessing RED zone. Confirmation must be stored and re-validated on session start. Under-18 users must be blocked with a clear message.

Technical Approach:

1. Add `age_verified: boolean` + `date_of_birth: date` to `user_profiles`
2. Create `AgeVerificationModal.tsx` — DOB picker, stores to profile
3. Wrap `ColorZoneSelector.tsx` — intercept RED selection with gate check
4. Add API middleware: `/api/chat/route.ts` rejects RED requests if `!age_verified`
5. Add gate check to `policy-router.ts`: if `zone=RED && !user.age_verified` → force YELLOW

DB Changes:

```
ALTER TABLE user_profiles ADD COLUMN IF NOT EXISTS
  age_verified BOOLEAN DEFAULT FALSE,
  age_verified_at TIMESTAMPTZ,
  date_of_birth DATE;
```

UI Components:

- `AgeVerificationModal.tsx` — DOB picker with clear explanation, stores to Supabase
- `RedZoneGate.tsx` — wrapper that intercepts RED color selection

Acceptance Criteria:

- [] Unverified user clicking RED zone sees age gate modal
- [] DOB entered as <18 shows blocked message, does not unlock RED
- [] DOB entered as 18+ stores `age_verified: true` and allows RED access
- [] Server-side: `/api/chat` returns 403 if `zone=RED` and user not age-verified
- [] Age verification persists across sessions (stored in DB, not just `localStorage`)

Effort: 1 week

Owner: Blossom (backend gate) + Bubbles (modal UI)

TFR-004 — Zero-Retention Spec Clarification + Privacy Policy Correction

Priority: ● P0 — Launch Blocker

Current State:

- `src/lib/email/templates/magic-link.ts` says "Zero-Retention. Private."
- But `/api/memory/`, `/api/extract-memories/`, `/api/journey/memories/` all store to Supabase
- This is a direct contradiction and a potential GDPR violation (false representation)

Requirement:

All marketing copy, onboarding text, and email templates must accurately describe what data is stored. The "zero retention" claim must be scoped correctly or removed entirely.

Technical Approach:

1. Search all files for "Zero-Retention", "zero retention", "no data stored"
2. Replace with accurate copy:
"Your conversation routing is private and not logged.
Your conscious memories are stored with your consent and deletable any time."
3. Add "Delete My Data" button to user profile settings
4. Implement `/api/user/delete-data` route that cascades deletes across all user tables

Files to Change:

- `src/lib/email/templates/magic-link.ts` — remove false claim
- Any landing page or onboarding copy referencing zero-retention
- Privacy Policy (TFR-002) must explicitly describe memory storage

Acceptance Criteria:

- [] No file contains unqualified "zero-retention" or "no data stored" claims
- [] Privacy policy accurately describes: what is stored, where, for how long
- [] User can request full data deletion from profile settings
- [] Data deletion cascades across: memories, capsules, journal, profiles, subscriptions

Effort: 3 days

Owner: MO (copy) + Blossom (delete-data API)

TFR-005 — Smart-Home Control (Remove from Feature Claims or Implement Stub)

Priority: ● P0 — Launch Blocker

Current State:

- `src/lib/integrations/integration-registry.ts` lists `smart_home` as an integration type
- `src/data/action-types.ts` has smart home action categories
- Zero implementation: no Google Home, HomeKit, MQTT, or Z-Wave code

Requirement:

Either (A) implement a minimum viable smart-home integration (Google Home API or Home Assistant MQTT), OR (B) remove smart-home from all public feature lists and disable the integration registry entry until implemented.

Recommendation: Option B — remove from public claims at launch. Implement in P2.

Technical Approach (Option B — 1 day):

1. Set `smart_home` entry in `integration-registry.ts` to `available: false`
2. Add "Coming Soon" label to any UI that lists smart-home
3. Remove from landing page feature list
4. Add to post-launch roadmap

Technical Approach (Option A — 3 weeks, P2 scope):

1. Integrate Home Assistant REST API (open source, self-hosted, covers most smart devices)
2. Create `/api/integrations/smart-home/[action]` routes
3. Store Home Assistant URL + API token in BYO keys vault
4. Add device discovery, state read, and command execution

Acceptance Criteria (Option B):

- [] Smart-home not mentioned in any public-facing copy at launch
- [] UI shows "Coming Soon" badge, not an active control

Effort: 1 day (Option B) / 3 weeks (Option A)

Owner: Blossom

TFR-006 — CAP Orchestrator (Central AI Policy Gateway for Sub-Domains)

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- No `cap-orchestrator.ts`, no `/api/cap/` routes, no "World" abstraction
- Each sub-domain (Emergent, Social Army, Job Hunt, Journal) calls OpenAI directly — bypassing zone policy, safety guardrails, and keyword telemetry
- The flagship spec requires all Worlds to route through the CAP Orchestrator, inheriting the chosen backend

Requirement:

A unified AI gateway that all internal tools call instead of OpenAI directly. Must inherit the user's current zone (TEAL/RED/YELLOW), apply safety guardrails, route to the correct backend, and pipe keywords back to the Side Panel.

Technical Approach:


```

Week 1: src/lib/cap/orchestrator.ts
- class CAPOrchestrator
- accepts: { worldId, userId, messages, zoneHint?, tools? }
- reads active zone from user session
- passes through PolicyRouter (reuses existing policy-router.ts logic)
- returns { response, zone, modelUsed, keywords[] }

Week 2: /api/cap/route.ts
- POST /api/cap/route
- Auth required, rate-limited
- All internal tools POST here instead of /api/chat

Week 3: Migrate 3 sub-domains
- Emergent: replace direct OpenAI calls with /api/cap/route
- Social Army: caption generation → /api/cap/route
- Job Hunt: resume parsing + cover letter → /api/cap/route

Week 4: Keyword telemetry pipeline
- CAP response includes keywords[]
- Client receives keywords, updates Side Panel via localStorage + WebSocket push

```

API Endpoint:

```

// POST /api/cap/route
interface CAPRequest {
  worldId: 'emergent' | 'social-army' | 'job-hunt' | 'journal' | 'rgy'
  messages: ChatMessage[]
  zoneHint?: 'TEAL' | 'RED' | 'YELLOW'
  tools?: ToolDefinition[]
}
interface CAPResponse {
  content: string
  zone: string
  modelUsed: string
  tokensUsed: number
  keywords: string[] // for Side Panel telemetry
}

```

Acceptance Criteria:

- [] Emergent editor calls `/api/cap/route` for all AI completions
- [] Social Army caption generation calls `/api/cap/route`
- [] Self-harm pattern in any World → forces YELLOW zone (guardrail applies globally)
- [] Keywords from CAP responses appear in Side Panel within 2 seconds
- [] Failover works: primary model down → secondary used transparently

Effort: 4 weeks

Owner: Blossom (orchestrator + API) + Bubbles (Side Panel telemetry UI)

TFR-007 — Audio Cues (Wake Chime, Speak Ticks, DND Mode)

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- `src/lib/audio/audioContext.ts` — AudioContext manager for gesture-unlock; no sound generation

- `src/lib/audio/audio-score-service.ts` — ambient music oscillators; no UI cues
- Zero wake chime, zero speak ticks, no DND mode anywhere

Requirement:

The spec requires: Wake chime on activation, soft ticks at TTS start/stop, single neutral tick on error, plus volume/on-off/DND controls.

Technical Approach:

```
// Create: src/lib/audio/ui-cues.ts
// Uses Web Audio API — zero external files needed

export function playWakeChime(ctx: AudioContext) {
  // 220Hz → 440Hz sweep, 80ms, cosine ramp out
  const osc = ctx.createOscillator()
  const gain = ctx.createGain()
  osc.frequency.setValueAtTime(220, ctx.currentTime)
  osc.frequency.linearRampToValueAtTime(440, ctx.currentTime + 0.08)
  gain.gain.setValueAtTime(0.3, ctx.currentTime)
  gain.gain.exponentialRampToValueAtTime(0.001, ctx.currentTime + 0.08)
  osc.connect(gain); gain.connect(ctx.destination)
  osc.start(); osc.stop(ctx.currentTime + 0.08)
}

export function playSpeakTick(ctx: AudioContext) { /* 880Hz, 30ms */ }
export function playErrorTick(ctx: AudioContext) { /* 330Hz, 50ms */ }
export function setAudioDND(enabled: boolean) { /* localStorage */ }
export function setAudioVolume(level: 0 | 0.25 | 0.5 | 0.75 | 1.0) { /* gain node */ }
```

Call Sites:

- `src/hooks/useAIState.ts` — `playWakeChime()` on idle → listening transition
- `src/hooks/useElevenLabsTTS.ts` — `playSpeakTick()` on TTS stream start + end
- `src/app/api/chat/route.ts` error handler → `playErrorTick()` via client event
- `src/components/SettingsCube/SettingsCubeApp.tsx` — DND toggle + volume slider

Acceptance Criteria:

- [] Wake chime plays when user activates mic (not on page load)
- [] Soft tick plays at start of TTS audio; second tick at end
- [] Error tick plays on AI failure (not on user input errors)
- [] DND mode silences all cues; preference persists across sessions
- [] Volume control works independently of system volume
- [] `prefers-reduced-motion` users: audio cues still work (audio ≠ animation)

Effort: 2 weeks

Owner: Bubbles (Web Audio implementation + DND settings UI)

TFR-008 — Vocspad (Unified Type + Talk Input Component)

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- `src/components/chat/ChatContainer.tsx` — standard text input + separate mic button
- `src/app/api/stt/route.ts` — Whisper STT working
- No unified Vocspad component that merges text and voice in one surface

Requirement:

A single input surface where users can type and talk simultaneously. STT transcription overlays in the text field live. Mic toggle embedded inline. Graceful fallback to text-only without mic permission.


Technical Approach:

```
// Create: src/components/chat/Vocspad.tsx
interface VocspadProps {
  onSubmit: (text: string) => void
  placeholder?: string
  disabled?: boolean
}

// Internal state:
// - isListening: boolean
// - transcript: string (live STT overlay)
// - typedText: string (manual typing)
// - merged: typedText + transcript (displayed together)

// On mic toggle: call useSTT() hook → stream Whisper transcript into textarea
// On Enter/Submit: send merged text
// Graceful fallback: if navigator.mediaDevices unavailable → hide mic button, text-only mode
```

UI Spec:



[ ▶] Type or talk... [↵ Send]
[] Live transcription here...

-  icon pulses red when listening
- Live transcript appears in lighter color
- Manual typed text appears in full color
- Both merge on submit

Acceptance Criteria:

- [] User can type text and it appears in Vocspad
- [] User can click mic, speak, and live transcript appears as overlay
- [] User can mix typing and speaking in same input
- [] Submit sends combined text to AI
- [] If mic permission denied, text input works normally with no errors
- [] Works on mobile (touch-friendly mic toggle)

Effort: 2 weeks

Owner: Bubbles

TFR-009 — Special Moves: Wink, Trust Earned, Handoff (3 Missing)

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- `src/components/cube/AICuboidGLB.tsx` implements 5 of 8 Special Moves in animation switch
- Missing: 'Wink', 'Trust Earned', 'Handoff'

Requirement:

All 8 Special Moves defined in the flagship spec must be implemented. Each is UI-only, does not affect routing, and must respect timing tokens ($\leq 200\text{ms}$ swap, 150–300ms glow).

Technical Approach:

Add 3 cases to the `switch (currentMove)` block in `AICuboidGLB.tsx`:

```
case 'Wink':
  // YELLOW-only: quick lateral tilt + emissive blink, ~200ms
  if (color === 'YELLOW') {
    const winkProgress = Math.sin(moveTimerRef.current * 30)
    groupRef.current?.rotation.set(0, 0, winkProgress * 0.15)
    emissiveModifier = winkProgress > 0 ? 1.3 : 0.7
    if (moveTimerRef.current > 0.2) triggerMove(null) // auto-reset after 200ms
  }
  break

case 'Trust Earned':
  // Warm golden pulse + slow scale breathe – acknowledgement moment
  mat.emissive.lerp(new THREE.Color('#ffd700'), 0.08)
  innerRef.current.scale.setScalar(1.0 + Math.sin(moveTimerRef.current * 3) * 0.05)
  mat.emissiveIntensity = 0.6 + Math.sin(moveTimerRef.current * 2) * 0.2
  break

case 'Handoff':
  // Outward scale expansion + fade → signals transition to another agent/World
  const handoffScale = 1.0 + moveTimerRef.current * 0.4
  innerRef.current.scale.setScalar(handoffScale)
  if (mat.transparent) mat.opacity = Math.max(0.1, 1.0 - moveTimerRef.current * 0.6)
  if (moveTimerRef.current > 1.5) triggerMove(null) // auto-reset after fade
  break
```

Trigger Conditions:

- `Wink`: trigger in YELLOW zone after playful/humorous AI response
- `Trust Earned`: trigger after user completes first memory save or first RGY match
- `Handoff`: trigger when routing request to another agent or switching Worlds

Acceptance Criteria:

- [] All 3 moves render without console errors
- [] Wink only fires in YELLOW zone (guarded by color check)
- [] Trust Earned produces a visible golden pulse (not same as Resonance)
- [] Handoff produces visible scale-out + fade
- [] All respect `prefers-reduced-motion` (check in parent hook)
- [] Timing stays within $\leq 200\text{ms}$ for instantaneous moves

Effort: 1 week

Owner: Pushpa (animation) + Bubbles (trigger wiring)

TFR-010 — CQ↔CQ Connections Wired to Main Chat UI

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- `src/lib/cq-to-cq/` module exists: CQ number generator, WebSocket, WebRTC calls

- CQ numbers can be generated programmatically
- Not connected to any UI — users cannot discover their CQ number or connect to others

Requirement:

Users must be able to: see their CQ number in their profile, share it, receive connection requests from other CQ numbers, and initiate direct encrypted chat/call with a CQ contact.

Technical Approach:

1. Profile page: display CQ number prominently (large, copyable)
2. QR code generation for CQ number (qrcode.react)
3. "Connect by CQ" button → input modal for entering another user's CQ number
4. WebSocket-based connection request flow (already exists in cq-to-cq/)
5. Accepted connection → appears in Connections list (new sidebar section)
6. From connection: initiate text chat or WebRTC voice/video call
7. DB: cq_connections table to store accepted connections

DB Changes:

```
CREATE TABLE cq_connections (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  requester_cq TEXT NOT NULL,
  recipient_cq TEXT NOT NULL,
  status TEXT DEFAULT 'pending', -- 'pending' | 'accepted' | 'blocked'
  connected_at TIMESTAMPTZ,
  created_at TIMESTAMPTZ DEFAULT now(),
  UNIQUE(requester_cq, recipient_cq)
);
```

UI Components:

- CQNumberDisplay.tsx — large CQ number + copy button + QR code
- CQConnectModal.tsx — enter CQ number, send request
- CQConnectionsList.tsx — accepted connections with chat/call buttons

Acceptance Criteria:

- [] Every user has a CQ number visible in their profile
- [] User can copy CQ number to clipboard
- [] User can enter another CQ number and send a connection request
- [] Recipient sees notification of incoming connection request
- [] Accepted connection enables direct chat via existing WebSocket infrastructure
- [] WebRTC call button works (uses existing cq-to-cq/ WebRTC implementation)

Effort: 2 weeks

Owner: Blossom (backend + WebSocket) + Bubbles (UI)

TFR-011 — Spending Caps Persistence (In-Memory → Supabase)

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- src/lib/spending/spending-cap-service.ts uses in-memory Map for usage tracking
- Server restart clears all spending cap data — users can accidentally overspend
- No DB table for usage records

Requirement:

Spending cap usage must persist across server restarts. Caps must be enforced even after Vercel cold starts. Usage history must be queryable for analytics.

Technical Approach:

1. Create `spending_usage` DB table
2. Replace in-memory Map with Supabase read/write in `spending-cap-service.ts`
3. Add daily/monthly rollup via Supabase cron
4. Cache in Redis (Upstash) for hot-path performance (optional P2)

DB Changes:

```
CREATE TABLE spending_usage (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  user_id UUID REFERENCES auth.users(id) ON DELETE CASCADE,  
  model TEXT NOT NULL,  
  tokens_used INTEGER NOT NULL DEFAULT 0,  
  cost_usd NUMERIC(10,6) NOT NULL DEFAULT 0,  
  period_start TIMESTAMPTZ NOT NULL, -- start of billing window  
  period_end TIMESTAMPTZ NOT NULL,   -- end of billing window  
  created_at TIMESTAMPTZ DEFAULT now()  
);  
CREATE INDEX ON spending_usage(user_id, period_start);  
  
CREATE TABLE spending_caps (  
  user_id UUID PRIMARY KEY REFERENCES auth.users(id) ON DELETE CASCADE,  
  daily_limit_usd NUMERIC(10,2),  
  monthly_limit_usd NUMERIC(10,2),  
  alert_threshold_pct INTEGER DEFAULT 80, -- alert at 80% of cap  
  updated_at TIMESTAMPTZ DEFAULT now()  
);
```

Code Changes:

- `spending-cap-service.ts`: replace `usageMap.get/set` with Supabase upsert calls
- Add `getRemainingBudget(userId)` that reads from DB
- Add cap enforcement to `/api/chat/route.ts` pre-call check

Acceptance Criteria:

- [] Spending usage persists after server restart (verify by redeploying and checking DB)
- [] User hitting 80% of cap receives email/in-app alert
- [] User hitting 100% of cap gets 429 response with clear message
- [] Usage history queryable in admin dashboard
- [] Latency: cap check adds <50ms to chat request (DB read with index)

Effort: 1 week

Owner: Blossom + Guy (DBA)

TFR-012 — Analytics Funnel Expansion (3 Events → Full Funnel)

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- `trackEvent()` is called in only 3 places across 300+ files

- No visibility into: signup conversion, feature adoption, retention, churn points, upgrade triggers
- Cannot make data-driven decisions without this

Requirement:

Instrument all critical user journey events. Enable full funnel analysis from landing → signup → first AI message → memory save → subscription.

Events to Add (minimum viable set):

Event	Trigger	Properties
page_view	Every route change	path , referrer
signup_started	Signup form opened	method (email/google/github)
signup_completed	Auth callback success	method , plan
first_message_sent	First chat message	zone , inputType (text/voice)
voice_activated	Mic button clicked	zone
memory_saved	Memory extract confirmed	memoryType
byo_key_added	BYO key saved	provider
social_post_created	Social Army post queued	platform , contentType
journal_entry_created	Journal saved	colorZone , promptUsed
capsule_matched	RGY match found	matchScore , zone
cubikey_viewed	CubiKey page opened	currentPlan
checkout_started	Stripe checkout opened	plan , billingInterval
subscription_activated	Webhook: sub created	plan , mrr
feature_gate_hit	Upgrade prompt shown	featureName , currentPlan
session_ended	Tab close / 30min idle	sessionDuration , messageCount

Technical Approach:

```
// Extend src/lib/analytics/analytics.ts (or create if not exists)
// Use PostHog JS SDK (already in stack per Section 1.1)

import posthog from 'posthog-js'

export function track(event: string, properties?: Record<string, unknown>) {
  if (typeof window !== 'undefined') {
    posthog.capture(event, properties)
  }
}
```

Acceptance Criteria:

- [] All 15 events above fire in PostHog when triggered
- [] Signup funnel visible in PostHog Funnel analysis
- [] Feature gate events show which features drive upgrades

- [] Session duration trackable for retention analysis
- [] Events fire in production (not just development)

Effort: 1 week (spread across components)

Owner: Bubbles (client-side events) + Blossom (server-side events)

TFR-013 — Social Army Poster Review Gate

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- `src/lib/social-army/social-army-service.ts` posts directly to platforms on schedule
- No human review step before posts go live
- No approval workflow, no preview UI before post
- If AI generates off-brand or harmful content, it posts immediately

Requirement:

All AI-generated social posts must enter a review queue before being published. User must approve each post (or configure auto-approve for trusted templates). Emergency stop to halt all scheduled posts.

Technical Approach:

1. Add status column to social_posts table: 'pending_review' | 'approved' | 'posted' | 'rejected'
2. After AI generation → set status='pending_review' (not 'approved')
3. Create SocialArmyReviewQueue.tsx — shows pending posts with preview + approve/reject/edit
4. Only approved posts are dispatched by the scheduler
5. Add "Emergency Stop" toggle in Social Army settings → sets global_pause=true in DB
6. Optional: auto-approve mode for power users who trust their templates

DB Changes:

```
ALTER TABLE social_posts ADD COLUMN IF NOT EXISTS
  status TEXT DEFAULT 'pending_review',
  reviewed_at TIMESTAMPTZ,
  reviewer_notes TEXT,
  auto_approved BOOLEAN DEFAULT FALSE;
```

UI Components:

- SocialArmyReviewQueue.tsx — scrollable list of pending posts, platform icons, preview text/image
- PostPreviewCard.tsx — individual post with approve/reject/edit buttons
- EmergencyStopToggle.tsx — red button in Social Army settings

Acceptance Criteria:

- [] New AI-generated posts default to `status='pending_review'`
- [] Scheduler only dispatches posts where `status='approved'`
- [] User can see all pending posts in review queue UI
- [] Approve/reject updates status immediately
- [] Edit redirects to post editor (preserves platform/scheduling)
- [] Emergency stop halts all scheduled dispatches within 60 seconds
- [] User can enable auto-approve mode per persona (stored in persona settings)

Effort: 2 weeks

Owner: Blossom (scheduler gate) + Bubbles (review queue UI)

TFR-014 — Onboarding Flow with 5-Branch Persona Routing

Priority: 🟡 P1 — 30-Day Post-Launch

Current State:

- No structured onboarding flow exists
- New users land directly in the main chat after signup
- No persona selection, no feature walkthrough, no Aha moment orchestration

Requirement:

After signup, new users complete a 5-question onboarding that routes them to the relevant Cubiqo sub-domain and populates their initial settings.

Onboarding Flow Design:

Step 1: "What brings you to Cubiqo today?"

- | | | |
|------------------------------|---------------------|----------------------------|
| → [A] Run my business better | → Solopreneur track | → Social Army + Agents |
| → [B] Find a job | → Job Seeker track | → Job Hunt + Resume |
| → [C] Process my thoughts | → Wellness track | → Journal + RGY |
| → [D] Build something | → Developer track | → Emergent + BYO |
| → [E] Connect with people | → Social track | → RGY Matching + CQ Number |

Step 2: "How do you prefer to interact?"

- Voice / Text / Both

Step 3: "Do you want Cubiqo to remember things between sessions?"

- Yes (enable memory) / No (session-only)

Step 4 (conditional on track A/D): "Do you have your own AI API keys?"

- Yes → BYO setup prompt / No → continue with Cubiqo keys

Step 5: "Give me your CQ number — share it with people you trust."

- Display CQ number + QR code + copy button

Technical Approach:

1. Create /onboarding route with multi-step form (5 steps)
2. Store onboarding responses to user_profiles.onboarding_data (JSONB)
3. Based on track selected: pre-configure color zone, feature flags, first suggestion
4. Redirect to track-specific dashboard section after completion
5. Skip link for power users

DB Changes:

```
ALTER TABLE user_profiles ADD COLUMN IF NOT EXISTS
  onboarding_completed BOOLEAN DEFAULT FALSE,
  onboarding_track TEXT, -- 'solopreneur' | 'job-seeker' | 'wellness' | 'developer' | 'social'
  onboarding_data JSONB DEFAULT '{}';
```

Acceptance Criteria:

- [] New users redirected to /onboarding after first login

- [] 5 steps render with clear UX (progress bar, back navigation)
- [] Responses stored to `user_profiles.onboarding_data`
- [] Track selection configures initial color zone preference
- [] Memory preference stored and applied immediately
- [] CQ number displayed prominently in Step 5
- [] Skip link available (stores `onboarding_completed: true` immediately)
- [] Returning users bypass onboarding

Effort: 2 weeks

Owner: Bubbles (UI flow) + Blossom (data storage + feature flag routing)

TFR-015 — Calendar API Integration (Google + Outlook)

Priority: ● P2 — 90-Day Enhancement

Current State:

- `src/lib/channels/` has channel type definitions including calendar
- No Google Calendar or Microsoft Graph API calls exist
- The flagship spec lists "Calendar" as a feature

Requirement:

Users must be able to connect Google Calendar or Outlook Calendar and have Cubiqo read/create events on their behalf.

Technical Approach:

1. Add Google Calendar OAuth scope to Supabase Google provider config
2. Create `/api/calendar/events` GET/POST routes using `googleapis` SDK
3. For Outlook: use `@microsoft/microsoft-graph-client`
4. Store calendar tokens in `user_byo_keys` (reuse BYO vault infrastructure)
5. Add Calendar connection to `HandshakeWizard.tsx` setup flow
6. AI tools: `readCalendar()`, `createEvent()`, `findFreeSlot()`

Acceptance Criteria:

- [] User can connect Google Calendar via OAuth
- [] AI can read upcoming events: "What do I have tomorrow?"
- [] AI can create events: "Schedule a call with Sarah on Friday at 3pm"
- [] Events appear in Google Calendar after AI creation
- [] Token refresh handled transparently

Effort: 3 weeks

Owner: Blossom

TFR-016 — Wallet / Payments (DB Migration + Stripe Crypto)

Priority: ● P2 — 90-Day Enhancement

Current State:

- `src/lib/wallet/wallet-service.ts` implements QR-delayed-release escrow logic
- No DB migration for `payments` or `wallet` tables
- No Stripe (for fiat) or crypto gateway (MetaMask/Coinbase) wired

Requirement:

Users must be able to hold value in a Cubiqo wallet, make payments to other CQ numbers, and use QR-based delayed release for trust-based transactions.

DB Changes:

```
CREATE TABLE wallet_accounts (  
  user_id UUID PRIMARY KEY REFERENCES auth.users(id),  
  balance_usd NUMERIC(12,4) DEFAULT 0,  
  balance_crypto JSONB DEFAULT '{}', -- {ETH: 0.01, SOL: 0.5}  
  created_at TIMESTAMPTZ DEFAULT now()  
);  
  
CREATE TABLE wallet_transactions (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  sender_id UUID REFERENCES auth.users(id),  
  recipient_id UUID REFERENCES auth.users(id),  
  amount_usd NUMERIC(12,4),  
  currency TEXT DEFAULT 'USD',  
  type TEXT, -- 'send' | 'receive' | 'escrow' | 'release' | 'refund'  
  qr_code TEXT, -- for delayed release  
  released_at TIMESTAMPTZ,  
  created_at TIMESTAMPTZ DEFAULT now()  
);
```

Effort: 3 weeks

Owner: Blossom + Guy

TFR-017 — Referral Programme

Priority: ● P2 — 90-Day Enhancement

Current State:

- No referral code system, no referral tracking, no reward logic

Requirement:

Every user gets a unique referral link. Successful referral (referred user subscribes) earns the referrer 1 month free or commission credit.

Technical Approach:

1. Generate referral_code (8-char alphanumeric) on user creation
2. Append ?ref=CODE to invite links
3. Track ref code through signup → first subscription
4. Apply reward: extend subscription_period_end by 30 days
5. Use PartnerStack (see TFR-002 agency table) for affiliate tracking if scaling

DB Changes:

```
ALTER TABLE user_profiles ADD COLUMN IF NOT EXISTS  
  referral_code TEXT UNIQUE,  
  referred_by TEXT REFERENCES user_profiles(referral_code),  
  referral_reward_credited_at TIMESTAMPTZ;
```

Effort: 1 week
Owner: Blossom

TFR-018 — Geo-fence Radius UI in Side Panel / RGY Settings

Priority: ● P2 — 90-Day Enhancement

Current State:

- `capsule-manager.ts` and `discovery-service.ts` reference geofence in matching logic
- `rgy_capsules` DB table has `geofence_km` column (from migration audit)
- No UI exposes this — users cannot set their location radius preference

Requirement:

Users must be able to set their preferred matching radius (e.g., "within 50km") from the Side Panel or RGY settings.

Technical Approach:

- Add radius slider (5km / 25km / 50km / 100km / Global) to `ProMatchSettings.tsx` or `KeywordPanel.tsx`
- Store to `user_profiles.match_radius_km`
- Pass radius to `capsule-manager.ts` `findMatches()` call

Effort: 1 week
Owner: Bubbles

TFR-019 — SettingsCube Spoken Live Confirmation

Priority: ● P2 — 90-Day Enhancement

Current State:

- `SettingsCubeApp.tsx` shows text result after command execution
- No TTS spoken confirmation: when user says "lock to TEAL," cube changes color silently

Requirement:

After a successful SettingsCube command, Cubiqo must speak the confirmation back to the user (e.g., "Done — locked to Teal mode").

Technical Approach:

- After `executeCommand()` success in `SettingsCubeApp.tsx`, call `/api/tts/route.ts` with confirmation text
- Use the active color zone's voice tone for the confirmation

Effort: 3 days
Owner: Bubbles

TFR-020 — Fabric-Soft-Touch Cuboid Material Variant

Priority: ● P2 — 90-Day Enhancement

Current State:

- Glass ☒, Metal/Chrome ☒, Solid ☒ — all implemented
- Fabric/soft-touch material: not implemented (spec requirement)

Requirement:

A cuboid variant with a soft textile appearance — low metalness, high roughness, subtle normal map, warm diffuse color.

Technical Approach:

```
// Create: src/components/cube/FabricCube.tsx
// Three.js MeshStandardMaterial:
const mat = new THREE.MeshStandardMaterial({
  color: new THREE.Color('#e8ddd0'), // warm cream – can be tinted by zone color
  metalness: 0.02,
  roughness: 0.95,
  normalMap: fabricNormalTexture, // procedural or loaded texture
  normalScale: new THREE.Vector2(0.3, 0.3),
})
```

Effort: 1 week

Owner: Pushpa

Summary: Complete Pending Work by Priority

PENDING FEATURES – PRIORITY MATRIX			
	Feature	Effort	Owner
P0	TFR-001 Stripe Billing + CubiKey	3 weeks	Blossom + Bubbles
P0	TFR-002 ToS + Privacy Policy	3 days	MO + Bubbles
P0	TFR-003 Age Gate for RED Zone	1 week	Blossom + Bubbles
P0	TFR-004 Zero-Retention Copy Fix	3 days	MO + Blossom
P0	TFR-005 Smart-Home (disable or implement)	1 day	Blossom
P1	TFR-006 CAP Orchestrator	4 weeks	Blossom + Bubbles
P1	TFR-007 Audio Cues (wake/tick/DND)	2 weeks	Bubbles
P1	TFR-008 Vocspad Unified Input	2 weeks	Bubbles
P1	TFR-009 3 Missing Special Moves	1 week	Pushpa + Bubbles
P1	TFR-010 CQ→CQ UI Wiring	2 weeks	Blossom + Bubbles
P1	TFR-011 Spending Caps → Supabase	1 week	Blossom + Guy
P1	TFR-012 Analytics Funnel (3 → 15 events)	1 week	Bubbles + Blossom
P1	TFR-013 Social Army Review Gate	2 weeks	Blossom + Bubbles
P1	TFR-014 Onboarding 5-Branch Flow	2 weeks	Bubbles + Blossom
P2	TFR-015 Calendar API (Google + Outlook)	3 weeks	Blossom
P2	TFR-016 Wallet / Stripe Crypto	3 weeks	Blossom + Guy
P2	TFR-017 Referral Programme	1 week	Blossom
P2	TFR-018 Geo-fence Radius UI	1 week	Bubbles
P2	TFR-019 SettingsCube Spoken Confirmation	3 days	Bubbles
P2	TFR-020 Fabric-Soft-Touch Cube Material	1 week	Pushpa

TOTAL EFFORT:

- P0 (all 5): ~5.5 weeks – DO THESE BEFORE ANY PUBLIC LAUNCH
- P1 (all 9): ~15 weeks – Complete within 30 days post-soft-launch
- P2 (all 6): ~10 weeks – Complete within 90 days post-launch

Solo developer path: ~30 weeks total to full spec compliance

2-developer path: ~16 weeks total

Table of Contents

Executive Summary — Product readiness scorecard, 3 launch blockers, key metrics

20 Techno-Functional Requirements — Every pending feature with full technical spec, SQL, API, effort estimate

- [1. Feature & UI Audit — What Is Actually Built](#)
- [2. Legal, Protection & Business Readiness](#)
- [3. Market Research & Monetisation Strategy](#)
- [4. Investor Strategy & Traction Plan](#)
- [5. MO's Final Verdict — Success Probability](#)
- [6. Patent Opportunities](#)
- [7. Architecture: Current State vs Target State](#)

8. [Patent Technical Flow Diagrams](#)

Appendices:

- [Appendix A — Key Numbers at a Glance](#)
 - [Appendix B — 30-Day Priority Action Order](#)
 - [Appendix C — Document Index](#)
 - [Appendix D — 15 Deep-Dive Topics](#) (*Dashboard, SEO, Social Army, Emergent, Agents, Duo Mode, Job Hunt, Journal, RGY, Users, Milestones, Marketing, Personas, Tricks*)
 - [Appendix E — Extended Strategy + Flagship Gap Analysis](#) (*Tools, Domains, Landing, CQ Score, CQ Number, Commerce, Silver Cube, BYO Offline, Browser Vision, User Personas, Flagship Spec Gap*)
-
-

1. Feature & UI Audit

1A — What Is Built and Working (Production-Ready)

Feature	Route / File	Evidence
Auth — Magic Link via Supabase	/auth	Fully wired, session-aware middleware
Chat + Voice — multi-turn with memory	/chat	ElevenLabs TTS + Whisper STT integrated
3D Plasma Cube — 120K particle WebGL	PlasmaWaveField.tsx	Real-time morph, 4 AI states
Voice State Machine	VoiceStateIndicator.tsx	READY→LISTEN→THINK→SPEAK states
Voice Modulation (Madhyama Marg)	voice-modulation.ts	4-mood parameter system, stochastic variance
TTS / ElevenLabs Streaming	/api/tts	Security headers fixed (PR #185)
STT / OpenAI Whisper	/api/stt	Full transcription route
BYO Mode — client-side key encryption	src/lib/byo/	AES-256-GCM + PBKDF2, 100K iterations
AI Policy Router — 5 zones	policy-router.ts	YELLOW/GREEN/RED/TEAL/FREEDOM + crisis override
LLM Router — 7 providers	llm-router.ts	Anthropic/OpenAI/Groq/Google/OpenRouter/Mistral
Journal CRUD — 24h gating	/api/journal/entries	Entries + history + stats + summary routes
Journal History API	/api/journal/history	Pagination + search working
Journey Memory — vector search	/api/journey/	pgvector, consent-gated, similarity API
AI Agents System	/api/agents/	Multi-agent with tool use
Browser Automation	/api/browser/	Playwright-backed agent capability
RGY Capsule Matching	rgy_capsules_and_matching.sql	4-stage staged matching with geofence
ProMatch Discovery	discovery-service.ts	Vector cosine opportunity scoring
Feature Flags	/api/feature-flags/	DB-backed, % rollout, admin UI
Security Headers	middleware.ts	OWASP-compliant, camera/mic fixed

Feature	Route / File	Evidence
Audit Logging	<code>admin_audit_log</code> table	All admin actions tracked
Rate Limiting	per-route middleware	Per-session request tracking
Social Army Worker	<code>social-army/worker.ts</code>	Queue processing with content engine
Job Hunt	<code>/job-hunt</code>	AI-assisted job search flow
Multimodal	<code>/api/multimodal/</code>	Image + voice + text combined
Code Execution	<code>/api/code/</code>	Sandboxed code runner

Total working API routes: 157

Total database migrations: 44

Database coverage: RLS on every table, pgvector enabled

1B — What Has Critical Gaps (Yellow — Partial)

Feature	Gap	Risk Level
Spending Caps	<code>let spendingRecord</code> — in-memory only, resets on every Vercel deploy	💰 Financial
Analytics	Only 3 events tracked (<code>magic_link_click</code> , <code>auth_modal_open</code> , <code>auth_complete</code>)	📊 Cannot steer
Journal History UI	API exists and works; frontend page not connected to it	UX gap
Onboarding	Saves to <code>localStorage</code> only; no email sent, no DB record	Retention
Social Army Poster	<code>poster.ts</code> posts directly via Puppeteer — no human-review gate enforced in code	⚠️ Legal risk
Pricing Page	Static React component — no Stripe checkout behind it	Revenue gap

1C — What Is a Placeholder (Grey — UI Exists, No Backend)

Feature	UI Location	Reality
CubiKey Page	<code>/cubikey</code>	"Beta" badge + 4 bullet points. No API, no portal, no billing
CubiKey Spec	<code>CUBIKEY_SPEC.md</code>	Detailed specification. Zero implementation

1D — What Is Completely Missing (Red — Launch Blockers)

Missing Item	Category	Why Critical
Stripe npm package	Billing	Zero revenue possible without this
Terms of Service	Legal	Stripe requires it before payments; GDPR requires it for EU users
/terms live page	Legal	No route exists in Next.js app
/privacy live page	Legal	Markdown doc exists; no live route
Cookie consent banner	GDPR	Required before any tracking
Refund policy	Legal	Stripe merchant agreement requires this
Email drip sequence	Retention	Day 1/3/7/14/30 — not wired
Referral programme	Growth	No referral_code column in user_profiles
Status page	Trust	Not implemented
support@cubiqo.ai	Support	Not set up

1E — Full Feature vs. Stub Assessment

FULL END-TO-END (API + DB + UI working):

- ✓ Chat + Voice (ElevenLabs + Whisper)
- ✓ BYO Mode (client-side AES-256-GCM)
- ✓ Journal CRUD (entries + history API)
- ✓ RGY Matching (4-stage SQL + vector)
- ✓ AI Agents (multi-tool, multi-model)
- ✓ Feature Flags (DB-backed, % rollout)
- ✓ Journey Memory (pgvector + consent)
- ✓ Auth (Supabase magic link)

PARTIAL (API works, UI gap or critical flaw):

- Spending Caps (concept correct, implementation wrong)
- Analytics (3 events instead of minimum 10)
- Journal History UI (API works, page not connected)
- Social Army (functional but legally unsafe)

PLACEHOLDER ONLY:

- CubiKey (landing page only)
- Founders Pass OAuth (disabled route)

MISSING ENTIRELY:

- Stripe/Billing
- ToS + Privacy live pages
- Email drip
- Cookie consent
- Referral programme

2. Legal, Protection & Business Readiness

2A — Legal Status: What Must Be Fixed Before First Paying User

Requirement	Status	Risk if Ignored
Terms of Service	✗ Missing	Stripe account termination; EU/CA regulatory action
Refund Policy	✗ Missing	Stripe requires before live payments; chargeback vulnerability
Privacy Policy (live page)	✗ Missing	GDPR €20M / 4% global revenue fine
Cookie Consent Banner	✗ Missing	GDPR non-compliance from Day 1
Age Gate (13+)	✗ Missing	COPPA violation in US; GDPR under-13 rules
AI Disclaimer	✗ Missing	FTC "AI disclosure" guidelines
Social Army ToS Consent	✗ Missing	Platform ban risk for Commander-tier users
Crisis Line in Journal	✗ Missing	Mental health data → duty of care obligation
GDPR Data Export endpoint	✗ Missing	EU users have right to their data
GDPR Data Deletion endpoint	✗ Missing	EU users have right to erasure

2B — The 10 Required Terms of Service Clauses

1. AI Output Disclaimer	"AI-generated content may be inaccurate. Not a substitute for professional advice."
2. Voice Recording Disclosure	"Voice data is processed by ElevenLabs and OpenAI. Not stored on Cubiqo servers."
3. BYO API Key Liability Limitation	"User-provided API keys are encrypted client-side. Cubiqo bears no liability for third-party API costs incurred through BYO keys."
4. Social Army Terms of Service Compliance	"Users must comply with LinkedIn, Twitter/X, and Instagram terms. Account bans resulting from Social Army automation are the user's sole responsibility."
5. Mental Health Disclaimer	"Cubiqo is not a mental health service. Journal contents are not monitored by clinicians. If in crisis, contact 988 (US) / 1-833-456-4566 (CA)."
6. Data Retention	"User data retained for [X] days after account deletion. Journey Memory data deleted immediately on opt-out."
7. Limitation of Liability	"Cubiqo's maximum liability is limited to the amount paid in the 3 months preceding the claim."
8. Governing Law	"[Province/State] law governs. Disputes resolved by binding arbitration."
9. Subscription & Refund Policy	"Monthly subscriptions cancel at period end. No refunds for partial months. Lifetime plans: 30-day full refund, no refunds after 30 days."
10. Age Requirement	"Users must be 13+ (16+ for EU residents under GDPR). No accounts for minors."

2C — Minimum Insurance for a Solopreneur AI SaaS

Policy	What It Covers	Priority	Estimated Annual Cost (CAD)
Cyber Liability	Data breach costs, regulatory defence, user notification	 Critical	\$800-\$2,500
Tech E&O (Errors & Omissions)	Claims that software caused financial harm	 Critical	\$1,200-\$3,500
General Liability	Third-party bodily/property damage	 Recommended	\$500-\$1,200
Media & Content	IP infringement claims on AI-generated content	 Important	\$600-\$1,800

Policy	What It Covers	Priority	Estimated Annual Cost (CAD)
Total estimated annual			\$3,100-\$9,000

Canadian-specific note: Incorporate provincially (Ontario or BC) before signing any insurance policy. Sole proprietors have unlimited personal liability. Incorporation + Cyber + Tech E&O is the minimum viable legal structure for launching a paid AI product.

2D — Adaptation Strategy by User Segment

Segment	Entry Point	Core Value Prop	Adaptation Needed
Solopreneur / Freelancer	Voice + Journal	"Your AI OS — remembers everything, executes anything"	Emphasise BYO mode (privacy); onboarding focused on voice first message
Developer / Technical	CubiKey API	"Cheapest multi-model router with privacy guarantees"	Needs CubiKey portal (6-8 weeks to build)
Social Media Creator	Social Army Commander	"Automate content across all platforms"	Social Army review gate must be enforced in code first
Job Seeker	Job Hunt	"AI co-pilot for your career pivot"	Weakest UX; needs resume parsing + application tracker

3. Market Research & Monetisation Strategy

3A — The Five Overlapping Markets Cubiqo Sits In

Market	2025 Size	CAGR	Cubiqo's Share Thesis
AI Personal Assistant	\$8.3B	28.5%	BYO privacy moat vs walled-garden assistants
AI Productivity Tools	\$12.1B	31.2%	Voice-first for solopreneurs
AI Journaling / Mental Wellness	\$4.1B	22.8%	Rozana Journal with RGY emotional tracking
AI Developer APIs	\$6.8B	45.3%	CubiKey multi-model routing (future)
Social Media Automation	\$9.7B	18.9%	Social Army Commander tier
Cubiqo addressable (overlap)	~\$2.1B	~28%	Conservative serviceable market

3B — Competitor Revenue Reality

Product	Category	Est. ARR	Why Cubiqo Is Different
Notion AI	Productivity	~\$50M	No voice, no memory, no agents
Monica AI	Personal assistant	~\$8M	No BYO, no voice state machine
Otter.ai	Voice transcription	~\$40M	Single function, no companion
Day One	Journaling	~\$12M	No AI, no voice, no emotional tracking
Buffer / Hootsuite	Social scheduling	\$100M+	No AI generation, no voice
OpenAI ChatGPT Plus	General AI	>\$1B	No voice companion, no memory, no social
Cubiqo	AI Companion OS	Pre-revenue	Multi-modal + memory + BYO + social army

Cubiqo's genuine differentiators:

1. BYO mode with client-side AES-256-GCM encryption — unique in the market
2. Voice + memory + agents in a single companion OS — no direct equivalent
3. RGY emotional state matching — novel peer discovery mechanism
4. 4-state AI visual companion (plasma cube) — creates emotional attachment

3C — The Three Revenue Engines (Sequenced)

Engine 1 — B2C Freemium (Start Here)

Free Tier:	Pro Tier (\$29/mo):	Lifetime (\$399 once):
10 voice msgs/day	Unlimited voice	Everything Pro, forever
3 journal entries	Advanced journal	Priority support
1 agent	10 concurrent agents	Early access to features
BYO mode only	Hosted mode + BYO	Founder-level CubiKey access
RGY basic	RGY ProMatch	Commander Lite included

What \$5K MRR looks like in this engine:

- 172 Pro subscribers at \$29/month, OR
- 420 free + 43 Pro (10% conversion, 430 free users total), OR
- Any mix that reaches \$5,000 recurring monthly

Engine 2 — B2D API (Build After \$1K MRR)

CubiKey Tiers:		
Starter \$29/mo	Pro \$99/mo	Enterprise (custom)
500 req/mo	5,000 req/mo	Unlimited
3 models	All 7 models	SLA + support
No reasoning	Reasoning mode	Custom endpoints

CubiKey is 6-8 weeks from launch-ready. The portal (/cubikey) is currently a placeholder.

Engine 3 — B2B Social Army (Build After \$5K MRR)

Commander Tier: \$499/month
Content queue: unlimited posts
Platforms: LinkedIn, Twitter/X, Instagram, Facebook
Content generation: AI-driven, brand voice trained
Human review gate: required before publish (must be enforced)
GFX Toolz: graphics generation included

Prerequisite: The human-review gate must be enforced in `poster.ts` code before this tier opens.

3D — Monetisation Tactics by Conversion Lift

Tactic	Conversion Lift (Industry data)	Status
Usage counter visible in UI ("7/10 free messages")	20-30%	✗ Not built

Tactic	Conversion Lift (Industry data)	Status
Upgrade modal at 90% limit (not 100%)	20-30%	✗ Not built
Yearly plan option (2 months free)	30-40% churn reduction	✗ Not built
Email drip with upgrade CTA at Day 7	15-25%	✗ Not built
Referral programme (give \$10/get \$10)	20-35% of new signups	✗ Not built
Feature teaser in UI ("unlock with Pro")	10-15%	✗ Not built
Free trial (7-day Pro access, no card required)	12-18%	Could build

All seven highest-ROI conversion tactics are not yet implemented.

4. Investor Strategy & Traction Plan

4A — Current State vs Investor Expectations

Metric	What Investors Expect for Pre-Seed	Cubiqo Today	Gap
MRR	\$0-\$2K (idea stage)	\$0	No billing
Users	50-500 (waitlist or beta)	Unknown	No analytics
Day-30 retention	20%+	Unknown	3-event tracking
Time to Aha moment	< 10 min	Unknown	Not instrumented
Working product	Yes (demo-able)	Yes 	—
Technical co-founder	Yes	Yes 	—
Defensible tech	Nice to have	Strong (BYO, voice, RGY)	—
IP / Patents	Nice to have	4 opportunities identified	File provisionals

4B — Realistic Investor Path

Available Now (No Metrics Required)

1. Friends & Family / Angels
 - \$10K-\$50K at SAFE note (post-money cap \$500K-\$1M)
 - Grounds: working product + technical credibility + defensible niche
 - Ask: enough to cover 6 months operating + patent provisionals (\$12K)
2. Government Non-Dilutive (Canada)
 - NSERC AI Alliance grants: up to \$100K non-dilutive
 - Canada Digital Adoption Program: up to \$15K
 - SR&ED tax credit: 15-35% of R&D labour costs
 - Grounds: AI research, privacy tech, voice AI

Available After \$1K MRR

3. Accelerators

Accelerator	Investment	Equity	What You Need
Y Combinator	\$500K	7%	\$1K MRR + traction
CDL (Canada)	Non-dilutive	0%	AI/tech focus
Antler	\$200K	10–12%	Solo founder OK
Google 4S	\$200K credit	0%	AI product req
BDC VC	\$500K–\$1M	Negotiable	Canadian, revenue

4C — The Investor Pitch Narrative (3 Chapters)

Chapter 1 — The Problem

"Every solopreneur runs 10+ tools that don't talk to each other. They forget context between sessions. They can't automate without a team. And every AI product either holds their data hostage or costs a fortune."

Chapter 2 — The Solution

"Cubiqo is the first AI companion OS that remembers everything, acts on anything, and runs entirely on your own API keys. Voice-first. Memory-first. Privacy-first. One interface for work, journaling, social, and automation."

Chapter 3 — The Moat

"Three things no competitor has together:

- (1) BYO mode with client-side AES-256-GCM encryption — your keys never leave your device.*
- (2) A persistent emotional memory system that gets smarter with every conversation.*
- (3) The RGY peer matching system — the only AI that understands your emotional state and connects you with people in the same headspace."*

4D — Metrics Gates Before Raising (In Order)

Gate 1: Ship Stripe + ToS (Week 1-2)
↓ first dollar of revenue
Gate 2: Hit \$1K MRR (Month 1-2)
↓ proof of willingness to pay
Gate 3: Hit \$5K MRR (Month 3-4)
↓ pattern of growth (not just one-time spike)
Gate 4: Day-30 retention $\geq$ 25% (Month 2-3)
↓ product actually retains users
↓
NOW RAISE: Angel round \$75K–\$250K at \$1M–\$2M valuation

4E — The 7 User Traction Implementations

Implementation	Expected Impact	Timeline	Cost
Referral programme	20–35% of new signups from referral	3–4 days	Low
Waitlist + invite-only beta	300–1,000 waitlist signups	1 day	\$0
Email drip (Day 1/3/7/14/30)	20–30% Day-30 retention improvement	3 days	Low
Public metrics dashboard (/open)	Press coverage + trust signal	1 day	\$0
GitHub stars campaign	Developer credibility signal	1 day	\$0
Product Hunt launch	500–2,000 sign-ups in 24h	2 days prep	\$0
10 user interviews (before any marketing spend)	Direction clarity	2 weeks	\$0

5. MO's Final Verdict — Success Probability

5A — What the Code Proves (Genuine Strengths)

Strength	Evidence	Why It Matters
Security quality	OWASP Top 10 covered, RLS on all tables, AES-256-GCM	Better than most Series A startups
BYO architecture	Client-side PBKDF2+AES, never transmitted to server	Trust moat; unique in AI market
Voice UX	4-state machine + 4-mood modulation + plasma cube morph	Creates emotional attachment
Feature flags	DB-backed % rollout; used throughout codebase	Ship safely solo
RGY matching	4-stage SQL with Haversine + vector overlay	Novel algorithm with patent potential
Spending caps concept	<code>SPENDING_CAPS</code> constants, monthly reset logic	Architecture is right

5B — What the Code Exposes (Real Risks)

Risk	Code Evidence	Severity
Spending caps in-memory	<code>let spendingRecord: SpendingRecord = {...}</code>	🔴 CRITICAL — resets on deploy
Analytics blind spot	<code>events.ts</code> tracks only 3 event types	🔴 HIGH — cannot steer product
No billing	Zero <code>stripe</code> imports anywhere in codebase	🔴 CRITICAL — existential
CubiKey placeholder	<code>/cubikey/page.tsx</code> is 40 lines of static HTML	🔴 HIGH — Revenue Engine 2 blocked
Social Army no gate	<code>poster.ts</code> calls <code>browser.launch()</code> directly	🟡 HIGH — legal liability
17 disabled routes	<code>route.ts.disabled</code> pattern throughout API	🟡 MEDIUM — incomplete features
ToS missing	No <code>/terms</code> route in <code>src/app/</code>	🔴 CRITICAL — legal blocker

5C — Four-Scenario Probability Assessment

Scenario	What Changes	P(\$5K MRR in 12 months)	P(Angel Round in 18 months)
A — Ship as-is	Nothing	8%	3%
B — 30-day pre-launch	Stripe + ToS + analytics + spending caps to DB	35-45%	25%
C — B + single positioning	"Solopreneur AI OS" as the hero, everything else as upsell	55-65%	45%
D — C + one part-time hire	Growth marketer or frontend contractor (10h/week)	65-75%	60%

Industry benchmark: CB Insights 2024 — 74% of bootstrapped AI products without live billing fail to reach \$1K MRR within 12 months.

5D — The Three Things NOT to Do Before Fixing the Basics

- ❌ Do NOT run Product Hunt before Stripe is wired.
 - A product that can't take money is a demo, not a product.
 - Product Hunt is a one-shot. Burning it on a demo is permanent.
- ❌ Do NOT acquire users before analytics diagnose retention.
 - Without 10+ tracked events, you have no idea why users leave.
 - Marketing spend before retention = pouring water into a leaky bucket.
- ❌ Do NOT open Social Army Commander tier before review gate is in code.
 - poster.ts currently posts directly. First ban will be within weeks.
 - A \$499/month tier with account-ban liability is a legal crisis waiting.

5E — The Single Most Important Strategic Decision

Pick one primary user. Build billing and retention loops around them.

The strongest candidate: **Solopreneur AI OS**

- User: solo professional aged 28-45, runs everything themselves
- Hero feature: voice + journal + memory
- Aha moment: "It remembered what I told it 2 weeks ago"
- Upsell: agents → job hunt → social army (in that order)
- Why: Most emotionally differentiated, most achievable at current dev stage, and the voice+memory combination has no direct competitor

6. Patent Opportunities

6A — The Four Genuine Opportunities (>50% USPTO Approval)

#	Patent Title	Key Files	Approval Est.	Commercial Value
P1	Content-Adaptive TTS Parameter Modulation	voice-modulation.ts , /api/tts	62-70%	● Highest — every AI voice product is potential infringer
P2	AI-State-Driven 3D Particle Morphology	PlasmaWaveField.tsx , EnergyCubeScene.tsx	58-66%	● High — visual AI interaction design
P3	Hierarchical Capsule Peer Matching	rgy_capsules_and_matching.sql	55-64%	● Medium — peer discovery niche
P4	Crisis Escalation Routing Override	policy-router.ts	52-60%	● Medium — safety routing architecture


6B — What the Code Actually Claims (Verified)

Patent 1 — Voice Modulation Pipeline

```
Input text (AI-generated)
|
| ▼ Stage 1: Keyword classifier (O(n), no API call)
|   intimate_markers[] → score_intimate
|   candid_markers[]   → score_candid
|   sincere_markers[]  → score_sincere
|   argmax → mood: intimate | candid | sincere | neutral
|
| ▼ Stage 2: LLM confirmation (TWEAK – add before filing)
|   Only if maxScore < 2 (ambiguous)
|   Haiku call → confirms or overrides
|
| ▼ Stage 3: 4D Parameter Vector Lookup
|   VOICE MOODS table:
|   sincere → {stability:0.75, similarity:0.75, style:0.15, boost:on}
|   candid  → {stability:0.40, similarity:0.70, style:0.65, boost:on}
|   intimate → {stability:0.60, similarity:0.85, style:0.25, boost:off}
|   neutral  → {stability:0.65, similarity:0.75, style:0.30, boost:on}
|
| ▼ Stage 4: Stochastic Variance Injection (±5%)
|   for each param p: clamp(0,1, p + (random()-0.5)*0.05)
|   → Madhyama Marg: prevents robotic/uncanny repetition
|
| ▼ ElevenLabs Streaming TTS API
|   POST /v1/text-to-speech/{id}/stream
|   voice_settings: {perturbed vector}
|   → audio/mpeg stream to client
```

Prior art distinguished: ElevenLabs (manual params), AWS Polly (SSML tags in text), Affective computing (audio input, not text input)

Patent 3 — RGY Matching + Journal Feedback Loop (Most Novel)



Most patent-valuable element: The closed loop — AI companion observes emotional state via journal → automatically encodes it in peer discovery capsule → peers matched on inferred (not self-reported) state. No prior art describes this feedback mechanism.

6C — Four Features Explicitly NOT Recommended for Patent

Feature	Why Not	Approval Est.
BYO AES-256-GCM encryption	Textbook Web Crypto API — Bitwarden/1Password prior art	15-25%
CubiKey intent routing	RouteLLM (Berkeley 2024), Martian, OpenRouter all describe this	20-30%
EMA adaptive learning model	Standard ML pattern; Netflix/Spotify well-documented prior art	25-35%
pgvector memory retrieval	LangChain, Mem0, Pinecone examples are all prior art	15-20%

6D — Filing Cost Table

Filing	Attorney est.	USPTO fee	Total	Urgency
Provisional P1 (Voice)	\$1,000-\$3,000	\$320	~\$3,320	🔴 File before any demo
Provisional P2 (Cube)	\$1,000-\$3,000	\$320	~\$3,320	🔴 File before any demo
Provisional P3 (RGY)	\$1,000-\$3,000	\$320	~\$3,320	🟡 File before launch

Filing	Attorney est.	USPTO fee	Total	Urgency
Provisional P4 (Crisis)	\$1,000–\$2,000	\$320	~\$2,320	🟡 File before launch
All 4 provisionals			~\$12,280	—
Full utility (12 mo later)	\$32,000–\$56,000	\$3,200	~\$48,000	12-month deadline
PCT international (12 mo)	\$4,000	\$4,000	~\$8,000	Optional

⚠️ **Critical:** Public disclosure (Product Hunt, blog post, demo video) before filing a US provisional destroys international novelty rights permanently in most jurisdictions. The US has a 1-year grace period. Canada, EU, Japan do not.

7. Architecture: Current State vs Target State

7A — Current Architecture (Color-Coded Reality)

Legend: ● Working | ● Partial | ● Placeholder | ● Missing



7B — Pre-Launch Target Architecture (30-Day Sprint)

Priority colour: ● P0 (do first) | ● P1 (week 2) | ● P2 (week 3-4) | ● Keep as-is

P0 – Week 1 (Legal + Billing: no revenue without these):

- | | |
|------------------------------------|----------------------------|
| ● Terms of Service (/terms live) | ● Refund Policy |
| ● Privacy Policy (/privacy live) | ● Cookie Consent Banner |
| ● Stripe npm install | ● Stripe Checkout API |
| ● /api/webhooks/stripe | ● subscription_tier column |
| ● Feature gates per tier | ● Spending caps → Supabase |
| ● Social Army review gate in code | ● AI disclaimer badge |
| ● Journal crisis line (988/116123) | ● Age gate at signup |

P1 – Week 2 (Launch quality):

- | | |
|---------------------------------|------------------------------|
| ● Analytics: 10+ events tracked | ● Welcome email on signup |
| ● Email drip Day 1/3/7/14/30 | ● Upgrade modal at 90% |
| ● Usage counters visible in UI | ● Journal History UI wired |
| ● BYO setup guide inline | ● Status page |
| ● support@cubiqo.ai active | ● Social Army consent screen |

P2 – Week 3–4 (Pre-GA enhancement):

- | | |
|----------------------------|------------------------------|
| ● Referral programme | ● Public metrics /open |
| ● GDPR data export/delete | ● Billing portal in settings |
| ● Annual plan (\$290/year) | ● Streak system (journal) |

7C — Post-Launch Architecture (After \$1K MRR)

REVENUE ENGINE 2: CubiKey API Portal (6–8 weeks)

- | | |
|-------------------------------------|------------------------|
| Developer Portal (/cubikey rebuilt) | API Key Generation |
| Usage Dashboard | Interactive Playground |
| Billing: \$29/\$99/Enterprise | Intent Router LIVE |
| Usage Metering (requests/tokens) | API Documentation |

GROWTH ENGINE:

- | |
|--|
| Product Hunt Launch (after provisionals filed) |
| Hacker News Show HN (technical article) |
| A/B testing: upgrade modal 3 variants |
| Annual plan option → 40% churn reduction |
| NPS survey at Day 30 (target: +50) |

ENTERPRISE / INVESTOR:


- | |
|---|
| Social Army Commander: review gate verified + rate limited |
| Data Processing Agreement (for EU enterprise) |
| Investor Data Room: Docsend + metrics dashboard + cap table |
| Angel Round: \$75K–\$250K at \$1M–\$2M valuation |
| Apply: CDL / Antler / YC (post \$5K MRR) |

7D — 6-Month Execution Timeline

Month	Revenue Target	Key Milestone	Patent Status
Now → 30 days	\$0 → first payment	P0+P1 sprint complete	File all 4 provisionals
Month 1	\$1K MRR	Beta launch (soft) · CubiKey build starts	Provisionals filed
Month 2	\$3K MRR	Email drip running · Day-30 retention measured	—
Month 3	\$5K MRR	Product Hunt launch · Angel conversations start	—
Month 4	\$8K MRR	CubiKey API portal live · Developer signups	—
Month 5	\$12K MRR	First enterprise Commander account	—
Month 6	\$20K MRR	Seed round prep · Accelerator applications	Full utility apps filed

8. Patent Technical Flow Diagrams

8A — Patent 1: Voice Modulation System Flow



What is NOT in prior art: The automated 5-stage pipeline (text → mood → vector → variance → stream) without human intervention per utterance.

8B — Patent 3: RGY Capsule Matching + Journal Feedback Loop

CLOSED LOOP (Most Novel Element):

Journal Entries (last 7 days)

- ↓ COUNT by color_category (green/yellow/red)
- ↓ argmax(distribution)
- ↓ UPDATE rgy_capsules SET color = dominant ← NO USER INPUT
- ↓

Capsule automatically reflects user's current emotional state

MATCHING PIPELINE:

All active capsules in database

- ↓ STAGE 1: WHERE color = query.color (hard gate – eliminates 60-80%)
- ↓ STAGE 2: AND intent = query.intent (hard gate – eliminates 50-70% of remainder)
- ↓ STAGE 3: calculate_keyword_match() (integer overlap – cheap)
- ↓ STAGE 4: WHERE distance_km ≤ radius (Haversine – optional)
- ↓ STAGE 5: 1-(embedding <=> query_embedding) (cosine – expensive – runs last on small set)
- ↓

Final ranked list: colour+intent filtered, keyword scored, geo constrained, vector re-ranked

DATABASE CONSTRAINT (novel semantic encoding):

```
CONSTRAINT valid_intent_for_color CHECK (  
  (color = 'yellow' AND intent IS NULL) OR      ← yellow = openness, no direction  
  (color IN ('green', 'red') AND intent IS NOT NULL) ← green/red = specific intent  
)
```

8C — Patent 4: Crisis Escalation Router

User Message → PolicyRouter.route()

- [RUNS FIRST] SAFETY PRE-PROCESSOR ← NOVEL: before any routing logic
 - Pattern: /self-harm|suicide|hurt myself|end my life/i
 - ↓ MATCH?
 - YES → ATOMIC OVERRIDE (NOVEL):
 - 1. zone = 'YELLOW' (unconditional – overrides even FREEDOM/RED)
 - 2. systemPrompt += crisis-support directive
 - 3. logSafetyOverride() → safety_override_log table [TWEAK]
 - 4. Continue to zone routing (now guaranteed YELLOW)
 - NO → Continue to zone routing normally

- ZONE ROUTING (runs after safety check)

- YELLOW → Llama 3.3 70B → Qwen Turbo → Haiku
- GREEN → MiniMax → DeepSeek V3 → Gemini Pro → Opus
- RED → Mixtral 8x22B → Llama Uncensored
- REASON → DeepSeek R1 → Opus
- SEARCH → [GPT-4o + Claude + Gemini + DeepSeek] in parallel → composite

WHAT PRIOR ART DOES DIFFERENTLY:

- OpenAI moderation: blocks content – does NOT reroute model or inject prompt
- Character.ai: filters output – not a routing-layer override
- Crisis apps: single-purpose – not general AI router with embedded safety

Summary: Pre-Filing Code Changes Required (Before Attorney)

Change	File	Time	Patent
Add <code>detectVoiceMoodDualStage()</code> (LLM 2nd stage)	<code>voice-modulation.ts</code>	1 day	P1
Add <code>MORPH_SPEEDS</code> record + use in <code>useFrame</code>	<code>PlasmaWaveField.tsx</code>	0.5 days	P2
Add journal→capsule colour trigger (SQL)	new migration	2 days	P3
Add <code>safety_override_log</code> table + <code>logSafetyOverride()</code>	new migration + <code>policy-router.ts</code>	1 day	P4
 DO NOT publish anything	—	—	All

Appendix A — Key Numbers at a Glance

Metric	Value	Source
Working API routes	157	<code>find src/app/api -name "route.ts"</code>
Database migrations	44	<code>supabase/migrations/</code> directory
UI components	80	<code>src/components/</code> directory
Disabled routes	17	<code>*.disabled</code> pattern
TODO/stub files	101	<code>grep -r "TODO\ FIXME" src/</code>
Analytics events tracked	3	<code>src/lib/analytics/events.ts</code>
Stripe npm packages	0	Confirmed: no stripe anywhere
Patent opportunities (>50%)	4	This analysis
Pre-launch P0 items	11	Section 7B
All-in provisional filing cost	~\$12,280	Section 6D
12-month \$5K MRR probability (Scenario C)	55–65%	Section 5C

Appendix B — 30-Day Priority Action Order

DAY 1–2 (This week – legal cannot wait):

- ☐ Retain patent attorney → file Provisional #1 (Voice Modulation)
- ☐ Draft Terms of Service (10 clauses from Section 2B)
- ☐ Install stripe npm package: `npm install stripe`
- ☐ Move spending caps from in-memory to Supabase (1 day fix)
- ☐ Add /privacy live route (docs already exist)

DAY 3–5:

- ☐ File Provisional #2 (Particle Morphology)
- ☐ Build Stripe checkout (Pro \$29/mo + Commander \$499/mo + Lifetime \$399)
- ☐ Build Stripe webhook → `subscription_tier` column update
- ☐ Add cookie consent banner (GDPR Day 1)
- ☐ Add AI disclaimer badge to every AI response

DAY 6–10:

- ☐ File Provisional #3 (RGY Matching) + #4 (Crisis Router)
- ☐ Wire analytics → 10+ events (`user_signed_up`, `first_voice_message`, `first_journal_entry`, `upgrade_modal_shown`, `subscription_created`, etc.)
- ☐ Add usage counters to UI ("7/10 free voice messages")
- ☐ Add upgrade modal at 90% of limit
- ☐ Enforce feature gates per `subscription_tier` in API routes

DAY 11–21:

- ☐ Welcome email via Resend/Loops on signup
- ☐ Email drip sequence (Day 1/3/7/14/30)
- ☐ Connect Journal History UI to existing API
- ☐ Social Army review gate enforced in `poster.ts`
- ☐ Status page (Betteruptime free tier → `status.cubiqo.ai`)
- ☐ `support@cubiqo.ai` active

DAY 22–30:

- ☐ Referral programme (`referral_code` + \$10/\$10 give/get)
- ☐ GDPR data export/delete endpoints
- ☐ Annual plan (\$290/year = 2 months free)
- ☐ Billing portal in `/settings`
- ☐ Public metrics `/open` page
- ☐ LAUNCH: soft beta to waitlist

AFTER LAUNCH (when \$1K MRR achieved):

- ☐ Product Hunt launch (Tuesday 12:01 AM PST)
- ☐ Hacker News Show HN
- ☐ Begin CubiKey portal build (6-8 weeks)
- ☐ Start angel conversations with metrics deck

Appendix C — Document Index

Document	Lines	Description
CUBIQ0_COMPLETE_REPORT.md	This document	True single-file consolidation of all analyses
CUBIQ0_MASTER_REPORT.md	951	Original 8-section master report
CUBIQ0_APPENDIX_B.md	945	15 deep-dive topics (see Appendix D below)
CUBIQ0_APPENDIX_C.md	1239	10 extended strategy topics (see Appendix E below)
CUBIQ0_ARCHITECTURE_CURRENT.md	230	Color-coded Mermaid current system
CUBIQ0_ARCHITECTURE_ROADMAP.md	254	Pre-launch + post-launch Mermaid roadmaps
PATENT_OPPORTUNITIES.md	637	Patent analysis with claim skeletons + specs
PATENT_FLOW_DIAGRAMS.md	580	4 patent flows + Gantt + pre-filing checklist
PRODUCT_LAUNCH_READINESS.md	463	Launch checklist, legal, insurance
GROWTH_AND_INVESTOR_STRATEGY.md	573	Market research, investor path, traction plan
MO_FINAL_VERDICT.md	261	Success probability, industry benchmarks

Appendix D — 15 Deep-Dive Topics

Topic 1 — Dashboard, Admin & Control Room Status

All Admin / Dashboard Routes Found in Codebase

Route	Page	Functional?	Notes
/admin	Admin dashboard	● Partial	Fetches live agent stats + usage data; spending caps display in-memory values
/admin/analytics	Analytics	● Partial	API exists; no funnel charts, only agent activity metrics
/admin/users	User management	● Working	Lists users, view sessions per user
/admin/feature-flags	Feature flags	● Working	DB-backed, % rollout, enable/disable
/admin/social-army	Social Army Console	● Working	Campaign progress, queue view, persona groups
/admin/spending	Spending limits	● Partial	Reads in-memory caps — shows wrong values after redeploy
/admin/journal	Journal admin	● Working	Admin view of journal entries
/admin/journey	Journey memory	● Working	Feature flag + metrics view
/admin/monitoring	Activity monitor	● Working	Real-time activity feed
/admin/noc	Network Ops Centre	● Partial	System health display
/admin/security	Security alerts	● Working	Failed logins, security alerts
/admin/self-heal	Self-heal logs	● Working	Automated health checks
/admin/system-health	System health	● Working	Database + services status
/admin/experiments	A/B experiments	● Working	Create/manage experiments
/admin/gate	Access gate	● Partial	UI exists, wiring unclear
/admin/designs	Design assets	● Partial	Preview system
/admin/events	Events log	● Working	Track analytics events
/admin/settings	Admin settings	● Working	Configuration
/dashboard	User dashboard	● Partial	

Route	Page	Functional?	Notes
			Stats load from DB; TODO in code: <code>LAST_THREAT_SCAN = '2 hours ago'</code> is hard-coded
<code>/dashboard/analytics</code>	User analytics	● Stub	All metrics show 0 — no data wired
<code>/founders-dashboard</code>	Founders view	● Partial	Different from founderspass
<code>/founders-pass</code>	Founder portal	● Working	Actions, audit, flags, security, sites
<code>/founders-pass/integrations</code>	Ecosystem integrations	● Working	Integration health, list
<code>/founderspass/dashboard</code>	FoundersPass dashboard	● Working	Feature flag toggles, duo mode visible
<code>/founderspass/experiments</code>	A/B testing	● Working	
<code>/agent-portal</code>	Agent portal	● Partial	Agent management UI

What Needs to Be Built for the "Control Room" to Be Complete

Priority 1 – Fix broken values (1 week):

- Spending caps page: move to Supabase so values survive deploys
- Dashboard: replace hard-coded "2 hours ago" with real timestamp
- `/dashboard/analytics`: wire to actual analytics events (sign-up funnel)

Priority 2 – Heavy Analytics Work (3-4 weeks):

- User funnel chart: signup → first voice → journal → upgrade
- Retention cohort table: Day 1/7/30 by signup week
- Revenue dashboard: MRR, churn, LTV, new vs renewing
- Voice usage heatmap: which hours users are most active
- Journal streak tracking: consecutive days + void detection
- RGY matching analytics: matches made, rooms joined, connections formed
- Agent usage by type: which A1-A7 agents are used most
- Social Army campaign ROI: posts → clicks → signups

Priority 3 – CEO View (2 weeks):

- Single-page CEO summary: MRR, DAU, CAC, LTV, churn, NPS
- Real-time alerts: spending cap approaching, error spike, churn event
- Competitor traffic comparison panel (SimilarWeb API)

Topic 2 — SEO & AI SEO Status

What Exists in Code

SEO Asset	File	Status	Quality
Canonical URL	<code>layout.tsx</code>	✓ Set	<code>https://www.cubiqo.ai</code>

SEO Asset	File	Status	Quality
OG tags	layout.tsx	✔ Set	Title, description, image
Twitter card	layout.tsx	✔ Set	summary_large_image
JSON-LD: FAQPage	layout.tsx	✔ Set	6 Q&A pairs
JSON-LD: SoftwareApplication	layout.tsx	✔ Set	Category, OS, description
JSON-LD: Organization	layout.tsx	✔ Set	Name, URL, logo
sitemap.xml	public/sitemap.xml	⚠ Outdated	Only 3 URLs; dated 2025-02-04
robots.txt	public/robots.txt	✔ Good	Allows all crawlers + AI crawlers
AI crawler allowlist	robots.txt	✔ Excellent	GPTBot, Claude-Web, PerplexityBot, Anthropic
llms.txt	Referenced in robots.txt	✗ Missing	File doesn't exist yet
Page-level metadata	Most pages	⚠ Mixed	Some pages have it, many don't
Core Web Vitals	Unknown	? Untested	Not measured
Sitemap coverage	3 pages	✗ Missing	Needs 20+ pages

What Needs to Be Done (SEO Roadmap)

Week 1 – Foundation (Technical SEO):

- Update sitemap.xml with all public pages:
/ /pricing /journal /chat /job-hunt
/founders /cubikey /privacy /terms
/blog (once live)
- Create /public/llms.txt
What LLMs should know about Cubiqo:
 - Product name and purpose
 - Key differentiators
 - API documentation summary
 - Pricing information
- Add page-level metadata to every page in src/app
- Add og:image to all product pages (use the 3D cube as visual)
- Measure Core Web Vitals via PageSpeed Insights

Week 2 – Content SEO:

- Create /blog with 3 seed articles:
 1. "Why BYO API keys are the future of private AI" (targets: private AI, BYO AI)
 2. "Voice-first productivity for solopreneurs" (targets: AI productivity, voice assistant)
 3. "RGY: How color signals change social matching" (targets: AI social, intent-based)
- Each article embeds the 3D cube demo as a hook
- Internal linking: blog → /chat → /pricing → /cubikey

Week 3 – AI SEO:

- Answer engine optimisation (AEO): structure FAQ answers as direct answers to queries like "best private AI assistant" and "AI productivity tool for freelancers"
- Submit to Perplexity, You.com, and SearchGPT via their partner programs
- Ensure llms.txt is comprehensive (product, pricing, features, API docs)
- Monitor Perplexity/ChatGPT citation tracking with a manual search query log

Ongoing:

- Publish 2 blog posts per week – target long-tail queries
- Monitor rankings: "BYO AI assistant", "voice AI for solopreneurs", "AI journal app", "private AI with own API key"
- Submit each new page to Google Search Console on publish
- Track AI citation rate: how often does Perplexity cite cubiqo.ai?

Topic 3 — The WeChat Vision & Affiliate Strategy

What "WeChat for AI" Means for Cubiqo

WeChat became a super-app by owning the user's daily life across: messaging, payments, mini-programs, social feed, and commerce — all within one UI. For Cubiqo, the equivalent is:

CUBIQO SUPER-APP VISION:

Layer 1: Voice + Chat (done )

→ The "messaging" layer. The daily touch point.

Layer 2: Journal + Memory (done )

→ The "diary" layer. Creates emotional lock-in.

Layer 3: RGY Matching (done )

→ The "social" layer. Connects users to opportunities and people.

Layer 4: Social Army (done, needs gate )

→ The "broadcast" layer. Users grow their audience from inside Cubiqo.

Layer 5: Job Hunt (done but no browser automation live )

→ The "commerce" layer. Real economic value delivered.

Layer 6: CubiKey (6-8 weeks to build )

→ The "mini-programs" layer. Third parties build on Cubiqo.

Layer 7: Emergent Platform (partially built )

→ The "app builder" layer. Users create new tools inside Cubiqo.

Layer 8: Integrations / Commerce (Shopify, Printify, Telegram  wired)

→ The "payments & commerce" layer. Revenue flows through Cubiqo.

Affiliate Strategy

TIER 1 – Referral Affiliates (launch day):

- Standard user referral: give \$10 / get \$10 off next month
- Tracked via referral_code column (needs to be added to user_profiles)
- Attribution: utm_source=referral&ref={code} in all shared links

TIER 2 – Content Creator Affiliates (\$1K MRR milestone):

- 20% recurring commission on referred Pro subscriptions
- Custom landing page: cubiqo.ai/{affiliate-slug}
- Dashboard: real-time earnings + conversion tracking
- Ideal affiliates: productivity YouTubers, AI newsletter writers, LinkedIn coaches

TIER 3 – Integration Partners (post \$5K MRR):

- Shopify app store listing: users connect their Shopify store to Cubiqo agents
- Printify partnership: Social Army designs → Printify products automatically
- LinkedIn Premium affiliate link: job hunt users who upgrade via Cubiqo link
- Revenue share: 10-15% of transaction value for commerce integrations

TIER 4 – CubiKey Ecosystem Affiliates (post \$10K MRR):

- Developer affiliates: build a plugin that uses CubiKey API, earn 10% of that plugin's revenue
- Agency affiliates: agencies manage Commander accounts, earn 15% recurring
- This is the WeChat mini-program model applied to Cubiqo

Topic 4 — Social Army: POC Status & 10×10×10 Assessment

What Exists in Code (Confirmed by Test Suite)

CONFIRMED WORKING (test suite passes):

- ✓ 10 platforms configured: twitter, linkedin, instagram, tiktok, youtube, reddit, pinterest, threads, facebook, discord
- ✓ 10 accounts per platform (persona groups in admin console):
The Builders (20), Productivity Gurus (30), Philosophers (15), Visual Artists (20), Memelords (15) → total 100 configured
- ✓ GFXToolz: processVideo(), createProject(), login() all tested
- ✓ Content queue: Supabase table "content_queue"
- ✓ Campaign progress: real percentage from DB (posted / target)
- ✓ Admin console: /admin/social-army – live campaign tracking
- ✓ Social campaigns table: "social_campaigns" in DB

WHAT 10-10-10 MEANS vs REALITY:

- 10 platforms: ✓ Configured in platforms.json
- 10 accounts: ✓ Persona groups defined (100 total across 5 types)
- 10 min posts: ⚠ Queue exists; cron job at /api/cron/rgy-discovery exists
but no confirmed 10-minute scheduler configured on Vercel

WHAT IS BROKEN / MISSING:

- ✗ poster.ts posts directly (no human review gate in code)
- ✗ Browser session management for 10 simultaneous platform sessions unclear
- ✗ GFXToolz processVideo() mocked – real API key needed for production
- ✗ No video/photo from actual Cubiqo app being captured and posted
- ✗ LinkedIn: rate limit = 1 post/day per account; can't do 10-min intervals
- ✗ Instagram: no direct API for posting without Meta Business approval
- ✗ TikTok: no direct API posting; requires TikTok for Developers approval

What Needs to Be Done to Reach 10×10×10

Phase 1 – Legal Safety (must do first):

- ❑ Add human review gate to poster.ts
 - Before publish: route to admin queue, require approval click
- ❑ Add Social Army Terms consent screen for Commander-tier users
- ❑ Add per-platform rate limits that respect each platform's ToS:
 - Twitter: 50 posts/day per account (10-min interval = fine)
 - LinkedIn: 1 post/day per personal account; use company pages
 - Instagram: Use Meta Graph API (requires approval)
 - TikTok: Official API only (requires developer approval)
 - YouTube: Title/description posts via YouTube Data API v3
 - Reddit: 0.9 posts/min across subreddits

Phase 2 – Real Content (media pipeline):

- ❑ App screenshot automation: Puppeteer takes screenshots of actual Cubiqo UI
 - Watermarked → Added to content queue automatically
- ❑ GFXToolz: plug in real API key and test processVideo() end-to-end
- ❑ Video: short-form clip of plasma cube voice interaction
 - Auto-rendered → posted to TikTok/Reels/Shorts
- ❑ Content variety: 60% educational, 30% product demo, 10% meme/entertainment

Phase 3 – Scheduling (10-minute cycle):

- ❑ Vercel Cron: configure /api/cron content-post at */10 * * * *
(Vercel Pro required for < 1 hour cron intervals; hobby plan = hourly minimum)
- ❑ Queue priority: round-robin across platforms to avoid burst
- ❑ Session persistence: store browser session cookies in Supabase, restore on next cycle (Playwright persistent contexts)

Phase 4 – Analytics:

- ❑ Track: impressions, clicks, profile visits, conversions per post
- ❑ A/B test content types weekly
- ❑ Automated report to admin console every Monday

Topic 5 — Emergent Coding Platform Status

What Actually Exists in Code

Component	File	Status
Projects API (create/list)	/api/emergent/projects/	✅ Full CRUD with auth + validation
Workspaces API	/api/emergent/workspaces/	✅ Working
Files API	/api/emergent/files/	✅ Working
Terminal API	/api/emergent/terminal/	✅ Working
Secrets management	/api/emergent/secrets/	✅ Full CRUD + rotate
Deploy API	/api/emergent/deploy/	✅ Vercel deploy wired
Analytics API	/api/emergent/analytics/	✅ Working

Component	File	Status
Audit log API	/api/emergent/audit/	✓ Working
Org management	/api/emergent/orgs/	✓ With RBAC
Monaco Editor	LiveCoderPane.tsx	✓ Working — TypeScript/JS/HTML/CSS
Sandbox preview	LiveCoderPane.tsx	✓ CSP-sandboxed iframe
Security RBAC	emergent/security/rbac.ts	✓ Role-based access control
Audit logger	emergent/security/audit-logger.ts	✓ Working
Stacks supported	Schema	TypeScript, JavaScript, Python, Next.js, React, Vue, Svelte, Vanilla

Honest Assessment vs Replit / Vercel v0

WHAT EMERGENT DOES WELL:

- ✓ RBAC with organisations and permission levels
- ✓ Secrets management with rotation
- ✓ Audit logging on all operations
- ✓ Live Monaco editor with syntax highlighting
- ✓ Sandboxed preview (CSP-strict)
- ✓ Vercel deployment integration
- ✓ Project + workspace management

WHAT IS MISSING FOR PROFESSIONAL GRADE:

- ✗ No collaborative editing (multiple cursors / real-time sync)
- ✗ No git integration in the UI (commits, branches, PRs shown in editor)
- ✗ No package manager UI (npm install within the platform)
- ✗ No terminal output streaming (real-time build logs)
- ✗ No debugging tools (breakpoints, step-through)
- ✗ No template library (start from Next.js starter, not blank)
- ✗ No AI inline code completion (Copilot-style)
- ✗ No multi-file navigation sidebar visible
- ✗ Preview only handles vanilla JS – not full Next.js / React rendering
- ✗ No mobile preview mode

What to Build to Make It Professional Grade

Sprint 1 (2-3 weeks) – Core editor UX:

- ❑ File sidebar: tree view, create/rename/delete files
- ❑ Terminal streaming: SSE-based real-time output (build logs, npm output)
- ❑ Template library: 5 starters (Next.js blank, React app, Landing page, API, Blog)
- ❑ AI code completion: tab-triggered inline Haiku/Deepseek completions

Sprint 2 (3-4 weeks) – Git integration:

- ❑ Connect to GitHub OAuth (already partially implemented in admin/connections)
- ❑ Show commits timeline in sidebar
- ❑ One-click "open PR" from Emergent to GitHub
- ❑ Auto-commit every save (like Replit)

Sprint 3 (2-3 weeks) – Full preview + deploy:

- ❑ Next.js preview: spin up a lightweight Next.js process per workspace
- ❑ Custom domain assignment for each deployed project (via Vercel API)
- ❑ Environment variable UI mapped to Vercel project env vars
- ❑ One-click deploy button → shows deploy status in real time

Topic 6 — In-App Agents: Capabilities, Access & Gaps

All 7 Agents (Confirmed in `src/lib/engine/router.ts`)

ID	Name	Trigger Keywords	Capabilities	Access Path
A1	Henry (Orchestrator)	Fallback / coordination	Multi-agent task orchestration, session memory, tool delegation	Default agent — active at all times
A2	Developer	code, dev, build, fix, bug, deploy, implement, refactor, debug	Code generation, file operations, terminal commands, code review	Say "write code" / "build" in chat
A3	Writer	write, document, content, blog, copy, article, draft	Long-form writing, SEO copywriting, email drafts, README	Say "write" / "draft" in chat
A4	Tester	test, qa, quality, verify, validation, regression, coverage	Test writing, bug reports, validation scripts	Say "test" / "QA" in chat
A5	Marketing	market, social, campaign, growth, promote, SEO, audience	Content calendars, social copy, campaign briefs, SEO articles	Say "marketing" / "campaign" in chat
A6	Animator	animate, visual, 3D, design, motion, graphic, render	Animation scripts, 3D prompts, visual concepts, design briefs	Say "animate" / "design" in chat
A7	Business	outreach, customer, sales, email, lead, CRM, contact	Sales emails, outreach templates, CRM data, business intelligence	Say "outreach" / "sales" in chat

Agent Capabilities in Detail

ALL AGENTS SHARE:

- ✓ Persistent session memory (conversation history stored)
- ✓ Tool registry access (filtered by user permissions)
- ✓ Auto-compaction when context window fills (saves tokens)
- ✓ SOUL.md personality file (agents/A1/SOUL.md etc.)
- ✓ Workspace isolation (each agent has its own file workspace)
- ✓ Multi-concurrent task support (A1: up to 2 parallel tasks)
- ✓ Spawn capability: A1 can spawn A2-A7 for subtasks
- ✓ Session compaction: auto-summarises old history

TOOLS AVAILABLE TO AGENTS (from tool registry):

- ✓ File read/write operations
- ✓ Code execution (sandboxed)
- ✓ Browser automation (Playwright via /api/browser/)
- ✓ Web search (via LLM router with search mode)
- ✓ Email send (via Resend integration)
- ✓ Calendar (if user connects Google Calendar)
- ✓ GitHub (if user connects GitHub)
- ✓ Slack, Discord, Telegram (if user connects)

How Users Access Agents

Current access paths:

1. /agents page – list all agents, start new session
2. Chat: type a message containing trigger keywords → routes automatically
3. /agent-portal – agent management UI
4. Voice: speak trigger keywords → voice routes to correct agent

What needs to be built:

- Agent picker UI in chat: @henry, @dev, @writer inline mentions
- Agent card gallery: visual showcase of each agent's specialty
- Agent output preview: show A3's drafts in a document editor UI
- Agent workspace browser: see A2's files without leaving chat
- Agent collaboration: A2 writes code, A4 auto-tests it, A3 documents it
(multi-agent pipeline initiated from single user request)

Topic 7 — Duo Mode & Companion Mode

What Exists in Code

DUO MODE (confirmed):

File: `src/components/chat/DuoModeToggle.tsx` – UI component exists
Feature flag: `duo_mode` in `feature-gate-simple.ts`
Access: Founders only (`FOUNDER_ACCESS.duo_mode = true`, `USER_ACCESS.duo_mode = false`)
Description: "Proactive AI interjections – AI interjects with health/safety/tone advice"
UI: Toggle button in chat header with purple glow animation
Backend: The autopilot profile extraction runs in background
(`/api/autopilot/tasks`) while AI companion converses

COMPANION MODE (Sidekick):

Feature flag: `sidekick_mode` – enabled for founders
File: `founderspass/dashboard` page shows "Sidekick Companion" as feature
Description: "AI companion mode"
Backend: No dedicated sidekick route found – appears to be a variant of chat mode

COPE MODE:

Feature flag: `cope_mode` – enabled for founders
Description: Not documented in code, but exists as a feature gate
Likely: Emotional support / mental health focused conversation mode

VOICE MODE:

Feature flag: `voice_mode` – separate toggle
Working: `voice_mode` is the ElevenLabs TTS + Whisper STT pipeline

What Needs to Be Configured / Built

Duo Mode – what the proactive interjections should include:

- ☐ Health: "You've been talking for 45 minutes – take a breath"
- ☐ Safety: Crisis pattern detected (already in `PolicyRouter`) → proactive check-in
- ☐ Tone coaching: "That email sounds aggressive – want me to soften it?"
- ☐ Focus: "You've shifted topics 5 times – want to focus on one thing?"
- ☐ Finance: "You mentioned money stress 3 times this week – want to journal?"

Companion Mode – what makes it different from standard chat:

- ☐ Persistent persona: remembers user's name, preferred topics, past decisions
- ☐ Proactive memory: "Last week you said you wanted to learn Python – any progress?"
- ☐ Emotional check-in: opens sessions with "How's today going?"
- ☐ Voice-default: companion mode should default to voice, not text
- ☐ Daily greeting message: pushed at user's typical start time

What to build:

- ☐ Duo Mode: connect autopilot tasks to proactive interjection API route
 - Background task checks for health/focus/tone signals every N turns
 - Injects interjection into conversation as assistant message
- ☐ Companion Mode: separate system prompt that emphasises continuity + memory
- ☐ User settings panel: "Companion personality" slider (professional ↔ personal)

Topic 8 — Job Hunt: Ideal User Flow vs Current Reality

Ideal Complete User Flow

STEP 1 – INTAKE:

- User lands on /job-hunt → "Set up your profile"
- Upload resume (PDF/DOCX) ✓ exists at /api/job-hunt/resume
 - Enter LinkedIn URL ✓ field exists in JobHuntProfile type
 - Select: target roles, locations, work type, salary range ✓ all in schema
 - Answer questionnaire (common application questions: "Are you authorized to work?") ✓ /api/job-hunt/questions exists

STEP 2 – AUTOMATED SEARCH:

- Cubiqa browser agent goes to:
- LinkedIn Jobs ✓ (Playwright browser + credentials)
 - Indeed, Glassdoor, ZipRecruiter ✓ (platforms in JobPlatform type)
 - Company career pages (custom browser sessions) ✓ via browser-service.ts
- Filters jobs by profile preferences
Scores each job for match (AI scoring against resume)
Returns ranked list to user

STEP 3 – USER REVIEW:

- User sees feed of matched jobs with match score
- "Apply automatically" checkbox per job
 - "Customise cover letter" option
 - "Add to watchlist" (manual review later)

STEP 4 – AUTOMATED APPLICATION:

- For each approved job:
- Navigate to job URL via Playwright ✓
 - Fill application form fields from profile ✓ (credentials encrypted in DB)
 - Upload resume from profile ✓
 - Use stored questionnaire answers for common fields ✓
 - Submit → store application in job_applications table ✓

STEP 5 – TRACKING & ALERTS:

- User gets daily email report: applications submitted, responses received ✓
Interview detected → alert sent ✓ (interview_alert report type exists)
Status tracking: pending → applied → screening → interview → offer/rejected ✓

STEP 6 – INTERVIEW PREP:

- Not yet implemented – should be:
- AI researches company from website + LinkedIn
 - Generates likely interview questions
 - Voice practice: Cubiqa asks questions, user answers, feedback given

Current vs Ideal Reality

Step	Ideal	Current Code	Gap
Resume upload	PDF/DOCX/TXT	✓ Implemented	—
LinkedIn URL	Field exists	✓ In schema	Not validated
Target roles/skills	Full form	✓ Types defined	Setup UI at /job-hunt/setup
Questionnaire	Common questions	✓ API exists	Q&A content not seeded

Step	Ideal	Current Code	Gap
Job search (browser)	6+ platforms	🟡 Browser wired	Auto-search not scheduled
Job scoring	AI match score	❌ Not implemented	Needs LLM scoring step
User review UI	Ranked feed	🟡 Dashboard exists	Feed not built
Auto-apply	Browser fill/submit	🟡 Playwright ready	Submit flow not fully built
Email alerts	Daily + interview	✅ Report types exist	Email send not wired
Interview prep	Voice Q&A	❌ Not started	Future feature
Application tracking	Full status flow	✅ Types + API	Dashboard shows stats

What to Build to Complete Job Hunt

Week 1:

- ❑ Job search automation: schedule browser agent to search LinkedIn/Indeed using user's target roles + locations at configurable intervals
- ❑ Job scoring: LLM call to rate each job against resume (0-100 match score)
- ❑ Jobs feed UI: ranked list with match score, apply/skip buttons

Week 2:

- ❑ Auto-apply flow: complete the form-fill + submit sequence
- ❑ Wire daily email report via Resend
- ❑ Interview alert: detect "interview" in email subject → send notification

Week 3:

- ❑ Company research agent: before application, A7 researches company
→ cultural fit analysis → personalised cover letter
- ❑ Interview prep: A1 asks role-specific questions → user answers via voice
→ A1 gives feedback on content + communication

Topic 9 — Daily Journal: Full User Flow & What's Pending

Current Working Flow (Confirmed in Code)

Step 1: User goes to /journal

- App checks: has user journaled in the last 24 hours?
 - ✓ API: GET /api/journal?sessionId={id}&userId={uid}
 - ✓ Response: { canJournal: true/false, nextAvailable: ISO timestamp }
- If already journaled: shows "Come back tomorrow" gate with countdown
- If can journal: shows "Begin today's session"

Step 2: Journaling flow (/components/journal/JournalFlow.tsx)

- ✓ 8 guided prompts, BigBoss confessional style:
 1. "How are you feeling right now? Really feeling?"
 2. "What happened today that actually mattered?"
 3. "How did that make you feel?"
 4. "What color is today – Green/Yellow/Red/Orange?"
 5. "What did you learn?"
 6. "What's one thing you want to accomplish tomorrow?"
 7. "Big picture – where are you heading in the next month?"
 8. "Anything else you need to get off your chest?"
- ✓ User types response to each prompt
- ✓ "Next" advances through prompts
- ✓ Final prompt saves to database

Step 3: Save & Summary

- ✓ POST /api/journal/entries – creates entry with:
 - session_id, user_id
 - all prompt responses
 - color_category (Green/Yellow/Red from prompt 4)
 - duration_seconds
- ✓ Email summary queued for user after completion
- ✓ Analytics event tracking

Step 4: History

- ✓ GET /api/journal/history – paginated, searchable
- /journal/history page exists but is NOT connected to API
- ✓ GET /api/journal/stats – streak, total entries, color distribution
- ✓ GET /api/journal/summary – AI-generated weekly summary

What Is Pending for Full Journal Experience

Week 1 – Connect the UI:

- /journal/history page: wire to /api/journal/history
- Show streak counter: "🔥 7-day streak" in journal header
- Calendar heatmap: GitHub-style activity grid showing journal days
- Color distribution chart: what % of days were Green/Yellow/Red

Week 2 – Voice Journal:

- Voice mode in journal: user speaks responses → Whisper transcribes
- Cubiqo reads each prompt aloud via TTS before user responds
- Creates a 15-minute "audio therapy session" feel

Week 3 – Insights & AI:

- Weekly summary email: "This week you felt Yellow 4 times, Red 2 times..."
- Pattern detection: "You feel Red every Monday – notice that?"
- Connection to RGY: journal color auto-updates capsule (patent opportunity #3)
- Memory extraction: key facts from journal → added to Journey Memory

Week 4 – Social / Connection:

- RGY connection: if 3+ consecutive Red days → proactively suggest RGY match with someone in Green (who might be in a position to help)
- Crisis detection: if journal responses contain crisis patterns → immediately surface crisis resources + duo-mode check-in

Topic 10 — RGY Keywords & 3 Intent Types: DB Status

What Exists in Database (Confirmed)

```
-- rgy_capsules table (CONFIRMED in migration 20260218000200):
color    TEXT: 'green' | 'yellow' | 'red'    ← COLOR CONFIRMED
intent   TEXT: 'collaborate' | 'trade' | 'company' ← 3 INTENTS CONFIRMED
keywords JSONB[] ← max 50 keywords, GIN indexed ← KEYWORDS CONFIRMED

-- DB Constraint (semantically meaningful):
yellow → intent MUST be NULL (open/neutral state)
green  → intent MUST be 'collaborate' | 'trade' | 'company'
red    → intent MUST be 'collaborate' | 'trade' | 'company'

-- RGY Chat Rooms table (CONFIRMED):
name, color, intent, keywords, geofence, expiry, max_participants

-- API routes (CONFIRMED working):
POST /api/rgy/intents      ← register intent
GET  /api/rgy/opportunities/discover ← find matches
POST /api/rgy/opportunities/express-interest ← signal interest
GET  /api/rgy/subscription ← subscription status
```

The 3 Intents Explained

Intent	Meaning	Example Use Case	Matching Logic
collaborate	Creative/technical partnership		Match two GREEN/collaborate users who share keywords

Intent	Meaning	Example Use Case	Matching Logic
		"I want to build something with someone"	
trade	Exchange of services/goods	"I have X skill, need Y skill"	Match complementary keyword sets
company	Build a company together	"Looking for a co-founder"	Match users with complementary domains (tech + biz)

What Is Fully Functional vs Needs Work

FULLY FUNCTIONAL:

- ✓ capsule creation with color + intent + keywords
- ✓ 4-stage matching algorithm (color → intent → keyword → vector)
- ✓ RGYRooms component: browse rooms by color + intent
- ✓ Join room: user can enter a chat room
- ✓ Geofence filtering (Haversine SQL function)
- ✓ Vector cosine similarity (pgvector overlay)

NEEDS WORK:

- Keyword autocomplete UI: suggest keywords as user types
- Capsule editor UI: where does user create/edit their capsule?
(RGYColorSelector.tsx and KeywordPanel.tsx exist – need to be wired into a flow)
- Proactive Cubiqo matching: Cubiqo goes and finds matches without user searching
(see Topic 11)
- RealTime notifications: when a match is found → notify user
- Match quality feedback: "Was this a good match?" for ML improvement
- Journal → capsule color sync (the patent-critical closed loop)

Topic 11 — RGY Chat, Capsule Logic & Proactive Matching

Is the Full System Understood in the Codebase?

YES — the full system is designed and partially implemented.

CAPSULE = color × intent × keywords
(e.g., "GREEN × collaborate × [react, typescript, fintech]")

WHAT THE UI CONFIRMS:

- ✓ RGYColorSelector.tsx – color selection UI exists
- ✓ KeywordPanel.tsx – keyword entry UI exists
- ✓ RGYIntentKeywordList.tsx – intent + keyword display component
- ✓ RGYRooms.tsx – browse + join chat rooms
- ✓ OpportunityFeed.tsx – opportunity discovery feed component

WHAT THE DB CONFIRMS:

- ✓ rgy_capsules table with all fields
- ✓ rgy_chat_rooms table
- ✓ Staged matching functions in SQL
- ✓ user_intents + opportunities tables (with vector similarity)
- ✓ RGY chat from Emergent platform was pushed (/api/rgy/* routes exist)

Proactive Matching — What Cubiqo Should Do Without User Asking

PROACTIVE FLOW (to be built):

Step 1: Cubiqo reads the user's active capsule

→ color = GREEN, intent = collaborate, keywords = [AI, voice, nextjs]

Step 2: Cubiqo runs the matching algorithm against all active capsules

→ finds 3 potential matches (anonymised at this stage)

Step 3: Cubiqo presents options to user in chat:

"I found 3 people who might align with what you're building.
One is a designer (GREEN × collaborate × [figma, ui, saas]).
Want me to find out if they're interested?"

Step 4: If user says yes → Cubiqo sends anonymous signal to matched user

→ POST /api/rgy/opportunities/express-interest (with anonymous=true)

→ Matched user gets notification: "Someone with aligned interests found you"

Step 5: Both users consent to reveal → Chat room created

→ Room expires in 7 days if no activity

→ If both parties join → connection formed → may or may not reveal identity

WHAT IS ALREADY BUILT FOR THIS:

- ✓ /api/rgy/opportunities/discover endpoint
- ✓ /api/rgy/opportunities/express-interest endpoint
- ✓ Anonymisation: can be done at API layer before revealing user IDs
- ✓ Chat rooms: auto-created with expiry

WHAT NEEDS TO BE BUILT:

- Cron: /api/cron/rgy-discovery already exists! Schedule on Vercel
- Notification: when a match is found → push notification to user
- Consent flow: before identity reveal, get explicit consent from both
- Chat UI: integrate RGY room chat into main Cubiqo chat interface
- UI connection: make the RGY section visible from /chat sidebar

How to Kickstart RGY (Immediate Steps)

Week 1 — Make capsule creation easy:

- Add "Set your RGY Status" to onboarding flow (after first voice message)
- Simple 3-step capsule creation: pick color → pick intent → add keywords
- Show active capsule in chat header as a colored dot

Week 2 — Enable proactive matching:

- Activate /api/cron/rgy-discovery on Vercel Pro
- Add notification when match found: in-app banner + email
- Wire the consent flow before identity reveal

Week 3 — Show the feed:

- Add "RGY tab" in main navigation
- Show: your capsule | matches found today | rooms you're in
- Chat rooms: integrate into existing voice/chat interface

Topic 12 — Where to Get Test Users + First 1,000 Users

Getting Test Users (Before Launch)

PHASE 0 — Friends & Founders (0-50 users):

- ❑ Slack: personal network, founder communities
- ❑ WhatsApp: close circle + extended network
- ❑ Give them direct access URL + Founders Pass
- ❑ Conduct 10 structured user interviews (30 min each via Zoom)
Ask: What did they try first? Where did they get stuck? What surprised them?

PHASE 1 — Community Seeding (50-200 users):

- ❑ Reddit:
r/Entrepreneur, r/solopreneur, r/productivity, r/artificialintelligence
Post: "I built a voice-first AI OS that runs on YOUR API keys — here's a demo"
Not promotional — lead with a genuine experiment or finding
- ❑ Hacker News: Show HN post — leads with technical credibility
"Show HN: CubiQo — a voice AI companion with BYO API key encryption"
- ❑ Product Hunt upcoming page: collect upvote pledges before launch
- ❑ IndieHackers: post as a build-in-public story
- ❑ Typefully/Tweethunter: schedule 30-day Twitter/X thread series
Day 1: "I'm building a WeChat for AI..."
Day 5: "Here's why client-side AES-256 encryption matters for AI privacy"
Day 10: "Our voice AI detects your emotional tone and adjusts its voice parameters"
(Build-in-public creates organic followers who convert to users)

PHASE 2 — Creator Partnerships (200-500 users):

- ❑ Target YouTube channels: Ali Abdaal, Thomas Frank, Tiago Forte, Mike Peralta
→ Offer 6 months free Pro in exchange for honest video review
→ Audience: productivity-obsessed professionals (exact ICP)
- ❑ Target AI newsletters: TLDR AI, The Rundown, Superhuman
→ Offer exclusive beta access to their reader list
- ❑ Target LinkedIn creators: Lara Acosta, Justin Welsh (solopreneur audience)
→ Use Social Army to prepare personalised outreach for each

PHASE 3 — Paid Acquisition (500-1000 users):

- ❑ Start ONLY after Day-30 retention > 20% (otherwise you're wasting money)
- ❑ Channels: X/Twitter Ads (tech/AI audience), LinkedIn Ads (professional persona)
- ❑ Ad creative: 15-second screen recording of 3D cube morphing as user speaks
→ "This is your new AI OS. It runs on your API keys."
- ❑ Landing page A/B test: cube demo vs. problem statement vs. social proof

BEST PLATFORMS TO RECRUIT TESTERS:

- ❑ Betalist.com — pre-launch testers who specifically sign up for betas
- ❑ BetaFamily.com — engaged testers who give structured feedback
- ❑ TestFlight (iOS) / Google Play beta — for mobile future
- ❑ Slack communities: Online Geniuses, Product Hunt Makers, Indie Founders
- ❑ Discord servers: Midjourney, Notion, Obsidian communities

Topic 13 — Milestones Aligned to Social, Launch & Growth

The Full Timeline

Milestone	Target Date	Social Signal	Technical Gate
Patent provisionals filed	Day 7 from now	—	File before any public post
P0 sprint complete	Day 14	Soft tweet: "Billing is live"	Stripe + ToS + caps fixed
Waitlist open	Day 14	Twitter/LinkedIn: "We're accepting 100 beta testers"	/waitlist page
50 beta users	Week 3	Reddit + HN Show HN post	Feature flags set for beta
10 user interviews done	Week 4	Build-in-public post: learnings thread	Record + publish findings
Day-30 retention measured	Week 6	—	Analytics ≥10 events live
\$1K MRR	Month 2	Tweet: "First \$1K MRR 🎉"	Stripe verified
Product Hunt launch	Month 3	Full PH campaign	After provisionals + billing
Hacker News Show HN	Month 3	Technical deep-dive article	After PH launch
\$5K MRR	Month 4	LinkedIn post: "from 0 to \$5K MRR in 120 days"	CubiKey portal live
Angel conversations	Month 4-5	—	\$5K MRR + metrics deck
Social Army activated	Month 4	First content wave live	Review gate confirmed
First affiliate partner	Month 4	—	Referral programme live
Accelerator applications	Month 5	—	Metrics + traction proof
\$20K MRR	Month 6	"Crossing \$20K MRR 🚀"	Enterprise accounts
Seed round close	Month 9	—	Post-accelerator

Topic 14 — Who Are Cubiqo's Users

Primary User Persona: The Solopreneur Operator

PERSONA 1 – "The Solopreneur Operator" (highest value, highest fit)

Age: 27-42

Works: Freelancer, consultant, indie maker, content creator

Income: \$4,000-\$20,000/month (unpredictable)

Tools today: Notion, ChatGPT, Gmail, Calendly, LinkedIn, Stripe

Pain: "I run 10 apps and none of them talk to each other.

I forget context between sessions. I can't automate without a team."

What they'll pay for: \$29/month (2 hours of saved work at \$50/hr)

Aha moment: "It remembered what I told it 3 weeks ago"

Entry feature: Voice + Journal

Upsell path: Pro → Agents → Social Army

Secondary Personas

PERSONA 2 – "The Frustrated Developer"

Age: 24-35 | Works: Startup developer, side-project builder

Pain: "ChatGPT is too locked down. Claude is expensive. I want one interface."

Entry: CubiKey BYO mode

Aha moment: "I'm running 7 models through one UI for \$0.10"

Upsell: CubiKey API subscription for their projects

PERSONA 3 – "The Social Media Hustler"

Age: 20-32 | Works: Content creator, personal brand, online coach

Pain: "Posting across 10 platforms is exhausting. Hiring a VA is expensive."

Entry: Social Army (Commander tier)

Aha moment: "It posted 50 pieces of content this week without me touching it"

Upsell: Commander tier upgrade

PERSONA 4 – "The Job Seeker in Transition"

Age: 25-45 | Works: Career pivoter, tech professional, laid-off worker

Pain: "Job hunting is a full-time job. I'm applying to 50 places manually."

Entry: Job Hunt feature (free tier allows limited applications)

Aha moment: "It applied to 20 jobs last night while I slept"

Upsell: Pro tier for unlimited applications + priority alerts

PERSONA 5 – "The Emotionally Exhausted Professional"

Age: 28-50 | Works: High-pressure professional, remote worker

Pain: "I want to process my day but can't afford a therapist."

Entry: Journal (free, always accessible)

Aha moment: "It noticed I felt red 4 Mondays in a row and asked if I'm okay"

Upsell: Pro for voice journaling + companion mode

Topic 15 — MO's Insider Tricks to Make This Mainstream

These are not in the codebase. These are strategic insights from pattern recognition across comparable AI products.

The 7 Insider Moves

Trick 1: Lead with the Cube, Not the Features

The 3D plasma cube morphing as the user speaks is the most emotionally compelling thing about Cubiqo. Every marketing piece should open with a 10-second screen recording of the cube reacting to voice. No competitor has anything like it. This is your hook. The features explain later.

Trick 2: "It remembered what you told it 3 weeks ago"

The single most powerful thing that will make Cubiqo feel different from ChatGPT is persistent contextual memory. The Aha moment is when a user mentions something offhand, comes back 3 weeks later, and Cubiqo references it unprompted. Build the memory highlight explicitly — after each session, show the user "Here's what I learned about you today" in 1-3 bullets.

Trick 3: Make BYO Mode the LEAD, Not the Footnote

Every AI product is leaking user data to their servers. BYO mode with client-side AES-256-GCM encryption is a genuine differentiator. The message: "Your API key never leaves your device." Position this as the privacy-first choice. Privacy-conscious users are the most loyal and highest LTV. They will evangelize for you because it aligns with their identity.

Trick 4: The RGY Colour System as a Social Identity

People love colour-coded identity systems (Red/Blue politics, MBTI, Hogwarts houses). RGY gives users a new social identity: "I'm a GREEN collaborator in the AI space." Build shareable RGY capsule cards — a simple image that says "🟢 Collaborate × AI × Voice × NextJS" that users can post to LinkedIn. It creates curiosity about Cubiqo without explaining it.

Trick 5: Ride the "Agentic AI" Wave Before It Peaks

The market narrative in 2026 is shifting from "chatbots" to "agentic AI" — AI that actually does things, not just talks. Cubiqo's browser automation, job hunt, and Social Army are all agentic. The framing: "Cubiqo doesn't just answer questions. It sends emails, applies to jobs, and posts content while you sleep." This narrative is 6 months ahead of mainstream adoption. Get there before the crowded market does.

Trick 6: The \$0 CAC Flywheel

Every user who uses Social Army to post on LinkedIn/Twitter is also indirectly marketing Cubiqo. Build the Social Army content templates to include: "Made with Cubiqo" watermark (optional but incentivized with extra free posts). Estimate: if 100 Commander users each post 5 times/week across 3 platforms, that's 1,500 brand impressions/week from users, not your marketing budget. This is the WeChat flywheel — the product markets itself through use.

Trick 7: The "Voice Diary" Emotion Hack

Research shows that the most retained app categories are health trackers, note-taking apps, and social apps — because they become part of daily routine. The Journal feature, especially in voice mode, creates a daily ritual. The secret: after the first voice journal, send an email: "Your voice note from today is waiting. Want to hear what Cubiqo noticed?" Play back a 30-second AI summary of what the user said. This creates a loop: journal → playback → insights → journal again. This is more habit-forming than any gamification system.

Summary: What to Build Next (Prioritised by Strategic Impact)

Priority	What	Why	Who	Weeks
● P0	Stripe billing + ToS	Existential — no revenue without this	Blossom	1-2
● P0	File patent provisionals	Public launch destroys rights	Founder	1
● P0	Spending caps → DB	Financial risk in prod	Blossom	1
● P1	Analytics ≥ 10 events	Cannot steer without data	Blossom	2
● P1	Journal history UI connected	Primary retention feature	Bubbles	1
● P1	Duo Mode backend wired	Key differentiated feature	Blossom	2
● P1	RGY capsule creation in onboarding	Starts network effect	Bubbles	1
● P1	sitemap.xml updated + llms.txt	SEO / AI SEO	MO	0.5
● P2	Social Army review gate	Legal/reputational safety	Blossom	1
● P2	Job hunt search automation	Key upsell trigger	Blossom	3
● P2	RGY proactive matching cron	Network effect engine	Blossom	2
● P2	Emergent platform MVP	Launch tool for other products	Blossom	4
● Post	CubiKey portal	Revenue Engine 2	Blossom + Bubbles	8
● Post	Referral programme	\$0 CAC flywheel	Blossom	3
● Post	Voice journal mode	Retention + daily ritual	Blossom + Bubbles	3
● Post	Interview prep feature	Job Hunt completeness	Blossom	3
● Post	Blog (3 seed articles)	SEO + AI SEO	MO + A3	2
● Post	Social Army 10×10×10	Viral distribution	Blossom	4

Appendix E — Extended Strategy: 10 Topics + Flagship Spec Gap Analysis

Topic 1 — Tools for Success + 3rd-Party Agencies

1.1 Core Digital Tools Stack

Category	Tool	Purpose	Monthly Cost	Priority
SEO & Keyword Research	SEMrush	Keyword tracking, backlink audit, competitor research, content gap	~\$120/mo (Pro)	P0
Web Analytics	Google Analytics 4	Page views, user journeys, conversion funnels, audience segments	Free	P0
Session Recording	Microsoft Clarity	Heatmaps, session replays, rage-click detection	Free	P0
Product Analytics	PostHog (self-hosted)	Feature flags, funnel analytics, cohort analysis, A/B testing	Free (OSS)	P0
Event Pipeline	Segment.io	Route analytics events to GA4 + PostHog + email tools from one SDK	Free up to 1K MAU	P1
Email Marketing	Resend + React Email	Transactional + drip sequences (already partially wired in codebase)	~\$20/mo	P0
CRM & Lifecycle	HubSpot Starter	Lead capture from landing page, lifecycle stages, deal pipeline	~\$20/mo	P1
Social Scheduling	Buffer or Publer	Schedule posts from Social Army output across 10 platforms	~\$15/mo	P1
Video Creation	CapCut + ElevenLabs	Short-form video content; AI voiceover matching Cubiqo's voice engine	~\$22/mo	P1
Heatmaps (paid)	Hotjar	More powerful session analytics once user base grows	~\$32/mo	P2
A/B Testing	Vercel Edge Config + PostHog flags	Landing page variant tests (plasma vs silver vs tech wireframe)	Included	P1
Uptime Monitoring	BetterStack (Uptime)	Alert on downtime, track 99.9% SLA	Free tier sufficient	P0

Category	Tool	Purpose	Monthly Cost	Priority
Error Tracking	Sentry	Catch runtime errors before users report them	Free up to 5K events	P0
AI SEO Content	Surfer SEO	Optimize blog posts and landing copy for search intent clusters	~\$89/mo	P2
Link-in-bio	Linktree Pro	Central hub for all social profiles pointing to cubiqo.com/join	~\$9/mo	P1

Minimum viable stack cost (P0 tools only): ~\$140/month.

1.2 3rd-Party Agencies & Partners

A. Affiliate Platform (Enable Revenue-Sharing)

Platform	Why	Revenue Model	Setup Time
Impact.com	Industry gold standard; 75K+ brands; supports SaaS payouts	% of MRR attributed to referral	2-4 weeks
ShareASale	Cheaper, simpler; good for digital products	Flat CPA or % commission	1 week
PartnerStack	Purpose-built for SaaS; native Stripe integration	Tiered commission based on MRR	1-2 weeks

Recommendation: Start with **PartnerStack** (SaaS-native, Stripe-native, lowest friction). Graduate to Impact once you have 50+ active affiliates.

What this unlocks: Users who refer others get CubiKey credits or cash. Creators who embed your widget get 20-30% recurring commission. This is the WeChat mini-program parallel.

B. Performance Marketing Agency

You don't need a full agency yet. Use these instead:

Resource	What For	Cost
Mayple.com	AI-matched freelance media buyers; pay per performance	Project-based
Fiverr Pro (verified)	Meta Ads / TikTok Ads setup for a specific campaign	\$200-500 per campaign
GrowthMentor.com	1:1 session with growth marketer who's scaled a similar product	\$50-100/hr

C. Investor Access

Service	What It Does	Cost
	List your raise; angels invest in syndicates of \$1K-10K each	

Service	What It Does	Cost
AngelList Syndicates		5% carry
Wefunder	Regulation CF crowdfunding; get 500+ micro-investors from your own community	3-5% fee
Gust.com	Submit to accelerators and angel networks en masse	Free
Republic.co	Consumer-facing investment raises; good for brands with a community story	6% fee
DocSend	Share pitch deck + track who read it, how long per slide	~\$45/mo

Solopreneur reality: Your first \$150K-250K will come from your personal network + one or two angels who believe in you, not agencies. Use Wefunder to run a community round once you hit 500 active users — this turns users into stakeholders and creates viral word-of-mouth.

D. Legal & Compliance

Service	What For	Cost
Clerky.com	Delaware C-Corp formation, cap table, SAFEs	~\$800 one-time
Stripe Atlas	Quick incorporation + Stripe account in one flow	\$500 one-time
Termly.io	GDPR/CCPA compliant ToS and Privacy Policy generator	~\$30/mo
WithKawazu.com	AI-assisted software IP attorney (startup rates)	~\$200/hr

Topic 2 — Domain Strategy: cubiqo / coqo / ciqo + 30+ Country TLDs

2.1 What You Have (Assumed Asset Inventory)

Based on the brand:

Domain Set	Examples	Strategic Value
Core Brand	cubiqo.com, cubiqo.ai, cubiqo.io	Primary; highest authority
Phonetic alternates	coqo.com, ciqo.com, cubiqu.com	Brand protection; prevent typosquatting
Country ccTLDs	cubiqo.ca, cubiqo.co.uk, cubiqo.de, cubiqo.in, cubiqo.au, cubiqo.fr, etc.	Geo-targeting, local SEO, legal jurisdiction
New gTLDs	cubiqo.app, cubiqo.chat, cubiqo.tech	Product-specific sub-brands

2.2 Recommended Strategy (Priority Order)

TIER 1 – PROTECT (register immediately if not done)

cubiqo.com	← PRIMARY (all traffic lands here)
cubiqo.ai	← AI product credibility signal
cubiqo.io	← Developer/tech audience fallback
coqo.com	← Typo protection
ciqo.com	← Typo protection

TIER 2 – REDIRECT (buy, 301-redirect to cubiqo.com)

cubiqo.ca	→ cubiqo.com	(Canadian market, your home)
cubiqo.co.uk	→ cubiqo.com	(UK – biggest English market)
cubiqo.in	→ cubiqo.com	(India – massive AI market)
cubiqo.au	→ cubiqo.com	(Australia – Anglophone)
cubiqo.de	→ cubiqo.com	(Germany – GDPR-forward)
cubiqo.fr	→ cubiqo.com	(France – EU anchor)
cubiqo.app	→ cubiqo.com	(Mobile installs)

TIER 3 – FUTURE (hold, activate only at scale)

cubiqo.ae	– UAE launch (high AI adoption)
cubiqo.sa	– Saudi Arabia
cubiqo.sg	– Singapore (APAC hub)
cubiqo.br	– Brazil (fastest-growing AI market 2025)
cubiqo.mx	– Mexico (Latin America gateway)

2.3 Technical Implementation

Action	How	Impact
301 redirects	Vercel → Project Settings → Redirects, or Cloudflare Page Rules	Preserves link equity; all ccTLD authority flows to cubiqo.com
Hreflang tags	Add <code><link rel="alternate" hreflang="en-ca" href="..."></code> in <code>src/app/layout.tsx</code>	Signals Google to show correct URL to each country
Geo-targeting in GSC	Set cubiqo.com as international target in Google Search Console	Tells Google your site serves all markets
CDN edge locations	Vercel Edge Network handles this automatically	Serves from nearest edge to each user

2.4 The CQ/Coqo Brand Option

If you want a shorter consumer brand (like how Instagram → Meta), consider:

- **coqo.ai** as the consumer-facing social/RGY app brand
- **cubiqo.com** as the platform/enterprise brand
- This mirrors Google/YouTube, Facebook/Instagram, Salesforce/Slack

Cost to hold 30 domains: ~\$900/year. Worth every dollar for brand protection at this stage.

Topic 3 — Landing Page: Conversion Standards + Silver Wireframe Feedback

3.1 Industry Conversion Benchmarks

Landing Page Type	Median CVR	Top Quartile CVR	What "Convert" Means
SaaS free trial	2-5%	8-12%	Email sign-up
AI tool waitlist	8-20%	25-40%	Join waitlist
Consumer app (curiosity hook)	12-35%	40-60%	Tap to enter + complete onboarding
Product Hunt launch day	2-6%	10-15%	Upvote + sign-up

Your current landing does NOT have a clear CTA above the fold. "Tap to begin" only shows up after 3 seconds of animation. This is the single biggest conversion leak.

3.2 What the Industry Standard Landing Page Has

ABOVE THE FOLD (visible without scrolling, no animation required):

LOGO + wordmark (top left)

HERO HEADLINE — "What it does in 8 words"

SUB-HEADLINE — "Who it's for and what changes for them"

[PRIMARY CTA BUTTON] "Start Free" / "Join Waitlist"

Social proof: "12,000 people already using this"

VISUAL HOOK (cube / animation / screenshot)

BELOW THE FOLD:

Pain → Solution narrative

3 feature cards with icons

Testimonials / logos

FAQ

Second CTA

3.3 Silver Wireframe — MO's Design Feedback

The silver wire aesthetic is **the right direction**. Here is why and what to tweak:

Element	Current (Plasma Wave)	Silver Wire (New)	Verdict
Emotional tone	Mystical, spiritual	Precise, intelligent, premium	✅ Better for B2B + sophisticated users
Color temperature	Warm purple/cyan	Cool silver-chrome	✅ Signals "intelligence"
Animation speed	Fast plasma pulse	Slow elegant morph	✅ More premium feel

Element	Current (Plasma Wave)	Silver Wire (New)	Verdict
Background	Near-black + purple glow	Near-black + blue-silver glow	✓ Cleaner
Text legibility	White on dark — good	White-silver on dark — good	✓
CTA visibility	Hidden below fold	Still hidden — needs fix	✗ Fix this

Recommended tweaks to the silver cube landing:

1. **Add a CTA button overlay** — don't rely on "tap anywhere". Add:

[Start Free — No Card Required]

as a pill button below the tagline. Keep "tap to begin" as secondary text.
2. **Add a single social proof line** between the wordmark and the tagline:

Trusted by 1,200+ solopreneurs ← update this as your numbers grow
3. **Morph timing** — the current 4-second cube→sphere cycle is excellent. Keep it.
4. **Add a subtle scanline/grid overlay** at 3–5% opacity to reinforce the "intelligence dashboard" aesthetic.
5. **The corner decorative marks** (already in the code) are perfect — don't remove them. They signal precision.

A/B Test Plan:

```

Variant A: plasma-wave (current default)
Variant B: silver-wireframe (new)
Variant C: silver-wireframe + CTA button overlay

Metric: email capture rate on the landing → app entry
Tool: PostHog feature flags + Vercel URL params (?landing=silver-wireframe)
Duration: 14 days minimum, 500+ sessions per variant

```

Topic 4 — RGY Score / CQ Intelligence (CQ Score Design)

4.1 What Exists Today (Code Reality)

From `supabase/migrations/` and `src/lib/`:

- RGY zones are assigned per journal entry / chat intent
- Keywords are tracked with zone colour + intent type (collaborate/trade/connect)
- There is no "CQ Score" numeric value — only zone labels

4.2 The CQ Score Vision (Design Proposal)

CQ = Cubiqo Intelligence Quotient — a living, composite score that represents a user's **growth trajectory across all life dimensions**, not a judgment on where they are.

CQ SCORE ARCHITECTURE

CQ SCORE (100–999)

- DIMENSION 1: Clarity (RED zone activity)
"How much is actively changing in your life?"
Score inputs: journal frequency, topics flagged as high-change, emotional variance across entries
Range: 0–333
 - DIMENSION 2: Growth (YELLOW zone activity)
"How actively are you learning & connecting?"
Score inputs: new keywords, new intent types explored, RGY matches made, job hunt activity, courses taken
Range: 0–333
 - DIMENSION 3: Alignment (GREEN zone activity)
"How aligned are you with your intentions?"
Score inputs: journal consistency, completed goals, RGY connections that led to outcomes, positive sentiment
Range: 0–333
- BONUS (+1 to +99):
Connection bonus: friends added, messages exchanged
App engagement: streak days, features used

Key principle (from your brief):

"Red should not be derogatory — it is change, it is dynamism, it is life in motion."

The scoring system should **reward change**, not penalise it:

- A user deeply in RED (lots of change) should have a high CQ if they're journaling about it
- A user stuck in YELLOW forever with no growth should score lower than a dynamic RED user

4.3 ConstantQuestioner.com Integration

Domain tie-in: `constantquestioner.com` → the "CQ number" concept.

Concept	How It Works
CQ Score	Numerical 100–999, visible on profile
CQ Number	Already implemented (CQ-XXXX-XXXX format in <code>src/lib/cq-to-cq/cq-number-generator.ts</code>)
Tie the two	CQ Score is displayed alongside CQ Number on the user profile card
ConstantQuestioner	A content brand / quiz platform that feeds CQ Score inputs: complete a questionnaire, your answers adjust your CQ Score dimensions

Database changes needed:


```
-- Add CQ Score to user profiles
ALTER TABLE profiles ADD COLUMN cq_score INTEGER DEFAULT 100;
ALTER TABLE profiles ADD COLUMN cq_clarity INTEGER DEFAULT 0;
ALTER TABLE profiles ADD COLUMN cq_growth INTEGER DEFAULT 0;
ALTER TABLE profiles ADD COLUMN cq_alignment INTEGER DEFAULT 0;
ALTER TABLE profiles ADD COLUMN cq_last_calculated TIMESTAMPTZ;

-- CQ Score history for charting
CREATE TABLE cq_score_history (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  user_id UUID REFERENCES auth.users(id) ON DELETE CASCADE,
  cq_score INTEGER NOT NULL,
  cq_clarity INTEGER,
  cq_growth INTEGER,
  cq_alignment INTEGER,
  calculated_at TIMESTAMPTZ DEFAULT NOW()
);
```

API endpoint needed: POST `/api/cq/recalculate` — triggered after each journal entry or RGY match event.

4.4 Why This Is Powerful

1. **Gamification without toxicity** — the score grows with engagement, never punishes users
2. **Social identity** — "I'm a CQ-847" becomes a badge, like credit scores but positive
3. **Investor metric** — average CQ score growth over time = proof of product value
4. **Monetization hook** — "CubiKey Pro users score 34% higher CQ on average" — upgrade incentive

Topic 5 — CQ Number as Universal Identity: Permanent vs Rotating + Adoption Strategy

5.1 What Already Exists

In `src/lib/cq-to-cq/cq-number-generator.ts`:

```
Format: CQ-XXXX-XXXX (e.g., CQ-8F3A-2K9B)
Rotation: every 30 days (privacy feature)
Entropy:  $32^8 = \sim 1$  trillion combinations
QR code: generateCQNumberQRData() already implemented
Shareable link: https://cubiqa.com/add/CQ-XXXX-XXXX already implemented
```

5.2 The Problem: Rotating $\neq$ Identity

A rotating number is **great for privacy** but **terrible for identity**. Phone numbers don't change. Instagram handles don't change. For CQ Number to work as a persistent identity anchor, you need a two-layer system:

LAYER 1 – Permanent CQ Handle (like @username)
 CQ-8F3A-2K9B ← permanent, assigned at account creation, never changes
 Used for: profile sharing, QR codes, "find me" links
 Stored as: immutable in profiles.cq_handle column

LAYER 2 – Rotating Privacy Token (already exists, keep it)
 CQ-2K9B-TEMP ← rotates every 30 days
 Used for: anonymous RGY matching, temporary connections
 Stored as: cq_numbers table (already implemented)

Code change needed: Add `cq_handle` column to profiles as immutable; generator runs once at registration.

5.3 Why Would People Want a CQ Number? (Adoption Psychology)

The same reason people want a LinkedIn URL, an Instagram handle, a Linktree. The answer is always: **social proof + access**.

Incentive	Mechanic	Analogy
Exclusivity	First 10,000 users get "Founder CQ" prefix: CQ-F-XXXX	Discord Nitro early badge
Network effect	"Add me on Cubiqo: CQ-8F3A-2K9B" becomes a social phrase	Cash App \$cashtag
Verification	CQ-Verified badge for users with 90-day+ accounts + CQ Score > 500	Twitter/X blue tick
Commerce	"Share your CQ Number, your friend gets 10% off their first CubiKey"	Referral codes
Professional	Add CQ Number to LinkedIn bio, email signature	Personal website URL
Business card	QR code of CQ Number on physical card → your Cubiqo profile	WeChat QR code
Emergency protocol	"In crisis? Share your CQ Number with someone you trust — they can see your zone"	ICE contact

The WeChat parallel: In China, you don't exchange phone numbers, you exchange WeChat QR codes. CQ Number is the same play — but AI-native.

5.4 Implementation Checklist

- Add `cq_handle` (permanent) to profiles table – migration needed
- Generate `cq_handle` at account creation, never change it
- Show CQ Number prominently on profile page (above the fold)
- Add "Copy CQ Number" button with one-tap clipboard copy
- Add QR code display on profile (qrcode.react library, already in many OSS stacks)
- Add "Find by CQ Number" search on friends/connection page
- Add CQ Number to the onboarding flow as a celebratory moment:
 - "You are CQ-8F3A-2K9B – share this with anyone to connect"
- Add CQ Number to email signature template sent after registration
- Create `/add/[cq-number]` public landing page (already referenced in generator)

Topic 6 — Cubiqo as Commerce Layer / WeChat Super-App

6.1 What Already Exists (Confirmed in Code)

`src/lib/deals/deals-service.ts` — A full Groupon-style deals engine already exists:

```
// Already implemented:
detectInterestCategories(text) // NLP keyword → deal category
hasDealIntent(text)           // "discount", "deal", "save" → trigger
fetchDeals({ category, maxResults }) // returns curated deal catalog
getContextualDeals(userMessage) // full pipeline: detect → fetch → return

// 10 deal categories already catalogued:
// food, travel, shopping, entertainment, beauty, fitness,
// electronics, education, services + more
```

The gap: All deals are hardcoded mock data with `url: '#'`. No real affiliate connection yet.

6.2 WeChat Super-App Layer (What to Build)

CUBIQO SUPER-APP LAYERS

LAYER 1 (EXISTS): AI Conversation + Memory

User talks to Cubiqo → gets smart responses → memory stored

LAYER 2 (EXISTS): RGY Journaling + Peer Matching

User journals → RGY zone classified → capsule created → matched

LAYER 3 (EXISTS - MOCK): Deals & Commerce

User mentions "cheap flights" → deals engine returns offers

NEXT: Connect to real affiliate feeds (Commission Junction, CJ Affiliate)

LAYER 4 (PARTIAL): CQ-to-CQ Messaging & Calls

Friends system, direct messages, WebRTC calls all coded

NEXT: Make this the primary social layer

LAYER 5 (NOT BUILT): Mini-Programs / CubiKey Ecosystem

Third parties build modules that live inside Cubiqo

User pays once, gets curated apps from trusted ecosystem

This IS the WeChat mini-program play

LAYER 6 (NOT BUILT): Payments & Wallet

CubiKey credits as in-app currency

Stripe → CubiKey top-up → spend on deals, mini-programs, upgrades

LAYER 7 (NOT BUILT): Local Services

Geo-aware deals + CQ-matched service providers in your city

"You mentioned you need a therapist → here are 3 CQ-verified therapists near you"

6.3 Affiliate Commerce: How to Actually Monetize the Deals Engine

Phase 1 — Week 1-2: Wire real affiliate feeds

Affiliate Network	Best For	Integration
Commission Junction (CJ)	Travel, electronics, retail	REST API + XML feed
Rakuten Advertising	Fashion, beauty, lifestyle	API
Amazon Associates	Electronics, books	Product Advertising API
Groupon Getaways API	Travel, experiences	REST (requires approval)
Honey / Capital One Shopping	Price comparison	Partner programme

Revenue per transaction: 3–15% commission depending on category. At 1,000 active users making 1 purchase/month averaging \$50 → **\$1,500–7,500/month passive income from deals alone.**

Phase 2 — Month 2: Add deal personalization

Connect deals to RGY zone + journal keywords. If user's GREEN zone keywords include "travel" and "wellness" → surface only travel and spa deals. This is the contextual moat that generic deal sites don't have.

Phase 3 — Month 3: CQ-Verified Local Professionals

Allow therapists, coaches, designers to list their services. Cubiqo takes 15% booking fee. User pays via Stripe. Service provider gets a CQ Badge. This is the LinkedIn-meets-Groupon-meets-Airbnb layer.

6.4 Discount Code Engine (What the User Described)

"it has a code for discount AND IT knows where exactly to buy and it knows the user for a good contextual recommendation"

This is **exactly** the deals engine but with two additions:

1. **Promo code field** in the deal object (add `promoCode?: string` to the Deal type in `src/lib/deals/types.ts`)
2. **Contextual matching** — already implemented via `detectInterestCategories()`, just needs real data

```
// Proposed enhanced Deal type
interface Deal {
  // ... existing fields ...
  promoCode?: string // "CUBIQ020" for 20% off
  promoExpiry?: string // ISO date string
  deepLink?: string // Direct link with affiliate tracking parameter
  affiliateId?: string // For attribution tracking
  isPersonalized?: boolean // Flag if this was surfaced by context engine
}
```

Topic 7 — Silver Wire Cube: Implementation + Visual Guide

7.1 What Was Built (This PR)

New file: `src/components/SilverWireLandingCube.tsx`

The component implements:

SILVER WIRE CUBE — TECHNICAL BREAKDOWN

GEOMETRY:

- BoxGeometry(1.6, 1.6, 1.6, 3, 3, 3)
 - subdivided for smooth morph
 - EdgesGeometry(BoxGeometry(1.6,1.6,1.6)) for crisp lines

MORPH ANIMATION:

- t = 0: Pure cube wireframe
- t = 0.5: Midpoint sphere-ish form
- t = 1: Full sphere (all vertices projected to unit sphere surface)
- Cycle: 4 seconds full cube→sphere→cube via $\sin(t * 0.5)$
- Implementation: per-frame Float32Array lerp between box and sphere positions

COLOR PALETTE:

- Edges: #D8E4F0 (blue-silver, 0.85 opacity)
- Fill mesh: #C8C8D0 (gunmetal silver, wireframe: true, opacity: 0.18)
- Corner spheres: #E8F0FF (near-white specular)
- Corner emissive: #A8C0FF (very subtle blue glow)
- Rings: #B0C8E8 (cool blue-silver, 0.28–0.35 opacity)
- Background: #080A0F (near-black, cooler than pure black)

LIGHTING:

- Ambient: 0.6 intensity, #D8E8FF tint
- Point 1: [4,4,4], intensity 1.8, white
- Point 2: [-4,-2,3], intensity 0.8, #A0B8D8
- Point 3: [0,-4,-3], intensity 0.5, #C8D8F0

EXTRAS:

- 3 axis rings (torus geometry) at 0.28–0.35 opacity
- 8 corner accent spheres (tiny, 0.028 radius)
- 4 corner L-bracket decorative marks (CSS, not WebGL)
- Horizontal rule gradient accent lines (CSS)

7.2 How to Preview It Right Now

Option A — URL parameter (no code change needed):

```
https://your-staging-url.vercel.app/?landing=silver-wireframe
```

Option B — Set as default:

```
# In your .env.local
NEXT_PUBLIC_LANDING_DEFAULT=silver-wireframe
```

Option C — Switch in `src/config/landing.ts`:

```
defaultVariant: 'silver-wireframe' // change from 'plasma-wave'
```

7.3 Design Comparison Table

Property	Plasma Wave (current)	Tech Wireframe	Silver Wire (new)
Dominant hue	Purple / cyan	Blue / pink / orange	Silver / chrome / blue-silver
Energy	High, kinetic	High, electric	Low, meditative, precise
User feeling	"Mystical AI"	"Sci-fi tech"	"Premium intelligence tool"
Brand fit	Early-stage wow	Hacker/developer	Enterprise / discerning user
Animation	Particle storm	Energy pulses	Geometric morph
Load weight	High (120K particles)	Medium (shaders)	Low (simple geometry)
Mobile performance	Medium	Medium	✔ Best (lightweight)
Conversion hypothesis	Wonder → curiosity	Interest → excitement	Trust → action

MO's recommendation: Run the silver wireframe as Variant B in A/B test. If it wins (which I expect it will for your core solopreneur demographic who values intelligence and precision), set it as default.

7.4 Next Enhancements (Post-Launch)

Enhancement	Complexity	Impact
Voice-reactive morph speed	Medium	Cube morphs faster when user speaks
RGY zone coloring	Low	Edges shift to Red/Yellow/Green based on user's current zone
CQ Score pulse	Low	Opacity pulses at rate proportional to user's CQ Score
Mouse parallax	Low	Cube tilts toward mouse cursor for depth illusion
Click ripple	Low	Tap sends shockwave through cube edges

Summary Priority Table

Topic	Key Action	Priority	Owner	Time
Tools	Set up GA4 + Clarity + PostHog + Sentry	P0	Blossom	2 days

Topic	Key Action	Priority	Owner	Time
Tools	SEMrush subscription + initial keyword audit	P0	JO + MO	1 day
Tools	Resend drip sequence (3-email welcome series)	P0	Blossom	3 days
Agencies	PartnerStack affiliate platform integration	P1	Blossom	1-2 weeks
Domains	301 redirects from all ccTLDs to cubiqo.com	P0	D2	1 day
Domains	Hreflang tags in layout.tsx	P1	Bubbles	0.5 day
Landing	CTA button overlay on silver wire landing	P0	Bubbles	1 day
Landing	A/B test: plasma-wave vs silver-wireframe	P1	Blossom+PostHog	3 days
CQ Score	DB migration for cq_score columns	P1	Guy	1 day
CQ Score	Recalculation API endpoint	P1	Blossom	2 days
CQ Score	Profile page CQ Score display	P1	Bubbles	1 day
CQ Number	Add permanent cq_handle to profiles	P0	Guy	0.5 day
CQ Number	Onboarding celebratory CQ reveal moment	P1	Bubbles	1 day
CQ Number	/add/[cq-number] public landing page	P1	Bubbles	1 day
Commerce	Wire real affiliate API (CJ or Rakuten)	P1	Blossom	1 week
Commerce	Add promoCode field to Deal type	P0	Blossom	0.5 day
Silver Cube	✅ Component built (this PR)	Done	MO	Done
Silver Cube	Add CTA button to silver landing	P0	Bubbles	1 day
Silver Cube	A/B test setup	P1	Blossom	2 days

Section 8 — BYO Mode: The Honest Reality, What True Offline Means, and the Road There

8.1 What BYO Mode Actually Does Today (vs What Was Assumed)

There is a common misconception worth clearing up first.

Dimension	What People Think BYO Does	What It Actually Does
API call destination	Goes directly from user's browser to Anthropic/OpenAI	User's browser → Cubiqo server (/api/chat) → Anthropic/OpenAI (with user's key)
Key storage	Key stays in the user's browser only	Key is AES-encrypted and stored in Supabase <code>profiles.byo_config</code> column
Cloud involvement	None	Cubiqo server is always in the middle — it decrypts and forwards
Cost to user	Billed to their API account	Yes — but still routes through Cubiqo infra
Privacy	Complete	Cubiqo server sees the decrypted key at decrypt-time in memory

Current BYO is a billing model, not a privacy model. The user pays their own cloud bill. Cubiqo still processes the request server-side. This is fine for Phase 1 and is honest — it's stated correctly in the UI ("encrypted before being sent to our backend"). But it is NOT the same as "no cloud call at all."

8.2 Three Tiers of "Bring Your Own" — From Now to True Offline

TIER 1 (DONE — live now)

Browser → Cubiqo Server → Anthropic/OpenAI
(user's encrypted key decrypted here)

- ✓ User pays own API bill
- ✗ Cubiqo server still in the middle
- ✗ Not private from Cubiqo infra
- ✗ NOT offline

TIER 2 (8–12 weeks — browser-side inference)

Browser → WebLLM (WASM/WebGPU, runs IN browser tab)
Model: Phi-3-mini-4k (2.3 GB download once)
or Gemma-2B-IT (1.8 GB)

- ✓ Zero server roundtrip for inference
- ✓ Truly private — model runs on user's GPU
- ✓ Works offline (after first load)
- ⚠ Requires WebGPU (Chrome 113+, Edge 113+, no iOS Safari)
- ⚠ Slower than cloud (3–8 tok/s vs 50+ tok/s)
- ✗ Actions (email, calendar) still need internet

TIER 3 (16–24 weeks — desktop + mobile native apps)

Tauri Desktop App → Bundled Ollama (llama3.2:3b)
One-click install, no terminal needed

- ✓ Full offline inference, no WASM limitations
- ✓ File system access for local journal, photos, data
- ✓ Can talk to local services (printer, local server)
- ✓ macOS / Windows / Linux
- ⚠ First download ~4 GB (app + model)
- React Native + llama.cpp (iOS/Android)
- ✓ Background inference, native camera, contacts
- ⚠ App Store approval process (2–4 weeks)

8.3 What Already Exists in the Codebase (Useful Foundations)

Asset	File	Status	Notes
Ollama integration	src/lib/ai/ollama.ts	✓ Working	Server-side only; points to localhost: 11434
AI router (Ollama-first)	src/lib/ai/router.ts	✓ Working	Routes 80% to Ollama, 20% to cloud
PWA manifest	public/manifest.json	✓ Done	Icons 192/512, display: standalone
Service worker	public/sw.js	✓ Working	Network-first + offline.html fallback
ServiceWorkerRegistration			

Asset	File	Status	Notes
	src/components/ServiceWorkerRegistration.tsx	✓ Working	Registers in production
PWA install prompt	src/components/PWAInstallPrompt.tsx	✓ Working	Android + iOS manual instructions
Offline fallback page	public/offline.html	✓ Done	Basic "you're offline" page
BYO key storage	src/lib/byo/byo-manager.ts	✓ Working	AES-encrypted in Supabase
BYO UI	src/components/byo/BYOSettings.tsx	✓ Working	WCAG 2.1 AA compliant

Key gap: The Ollama route currently only works when Cubiqo's server has Ollama running locally (localhost:11434). That is the server's localhost, not the user's machine. To make it truly client-side offline, the model must run in the browser (WebLLM) or in a native app (Tauri + Ollama).

8.4 Tier 2 Implementation Plan — Browser-Side Inference (WebLLM)

Technology: [MLC-AI WebLLM](#) — runs quantised LLMs in browser via WebGPU. No Python, no Ollama, no server.

Step-by-step build plan:

```

Week 1-2: Research + Proof of Concept
- Install: npm install @mlc-ai/web-llm
- Create src/lib/ai/webllm-client.ts
- Wrap WebLLM engine in a React context (useWebLLM hook)
- Gate behind feature flag: feature_flag = 'browser_inference'
- Test with Phi-3-mini-4k-instruct-q4f16_1 (smallest viable model)

Week 3-4: UX for model download
- First-launch download modal: "Download 2.3 GB AI model once, run forever offline"
- Progress bar using WebLLM's initProgressCallback
- Cache model in IndexedDB (WebLLM handles this automatically)
- Show estimated time remaining

Week 5-6: Feature integration
- Chat route: if user has browser_inference enabled AND model loaded → use WebLLM
- Offline detection: navigator.onLine → switch to WebLLM automatically
- Memory extraction: run locally (smaller prompt, acceptable quality)
- RGY classification: run locally (keyword-based, already works without cloud)

Week 7-8: Actions bridge
- Email/calendar/social still need internet → show "action queued" when offline
- Sync queue: store pending actions in IndexedDB → execute when back online
- Show user clearly: "Thinking offline ✓ | Will send email when connected ↑"

```

Browser compatibility gate (must show before offering):

```
const supportsWebGPU = async (): Promise<boolean> => {
  if (!navigator.gpu) return false;
  const adapter = await navigator.gpu.requestAdapter();
  return !!adapter;
};
```

Browser	WebGPU Support	Notes
Chrome 113+	✓ Full	Best experience
Edge 113+	✓ Full	Same Chromium engine
Firefox	⚠ Partial	Behind flag, unstable
Safari 17.4+	⚠ Partial	Metal backend, limited
iOS Safari	✗ None	No WebGPU as of Feb 2026
Android Chrome	✓ Most devices	Depends on GPU

Models to offer (in order of size vs quality):

Model	Size	Speed	Quality	Use Case
Phi-3-mini-4k (q4f16)	2.3 GB	6-10 tok/s	Good	Default choice
Gemma-2B-IT (q4f16)	1.8 GB	8-12 tok/s	OK	Smaller devices
Llama-3.2-3B (q4f16)	2.0 GB	7-11 tok/s	Good	Open weights
Phi-3.5-mini (q4f16)	2.5 GB	5-8 tok/s	Best	Power users

8.5 Tier 3 — Desktop App with Bundled Ollama (Tauri)

Why Tauri over Electron: Tauri uses the OS WebView (WebKit/EdgeWebView2) instead of bundling Chromium. App size is ~10 MB vs ~150 MB for Electron. Rust backend is faster and more memory-efficient.

Architecture:

```
Tauri Rust Core
- Spawn Ollama as a sidecar process
- File system access: journals, photos, local data
- System tray icon: "CubiQo running"
- Auto-updater (built-in Tauri updater)
- Biometric auth (TouchID / Windows Hello)

Next.js frontend (compiled to static)
- All existing React UI unchanged
- AI calls → localhost:11434 (Ollama sidecar)
- Storage → user's local files + SQLite
```

One-click installer experience:

- 1. User downloads CubiQo-1.0.0-setup.exe (Windows) or CubiQo-1.0.0.dmg (macOS)
 - 2. Installer places Ollama binary as a Tauri sidecar
 - 3. First launch: "Downloading your private AI brain... [██████████] 3.2 GB" (model download)
 - 4. Model cached at ~/.cubiqo/models/ — never re-downloads
 - 5. All AI inference: local. Internet only for actions (email, web search, social posts)
- Timeline:** 16 weeks from decision to App Store submission (macOS/Windows). iOS/Android React Native variant: separate 20-week track.

8.6 Immediate Next Steps (What to Do Now)

Priority	Action	Owner	Effort
P0	Add honest copy to BYO UI: "Your inference still routes through Cubiqo servers"	Bubbles	0.5 day
P0	Add "True Offline (Beta)" toggle behind feature flag (disabled by default)	Blossom	1 day
P1	WebLLM PoC: <code>npm install @mlc-ai/web-llm</code> , single chat endpoint test	Blossom	3 days
P1	Enhance service worker: cache chat-related static assets for offline splash	Blossom	1 day
P1	Offline action queue in IndexedDB for emails/posts when offline	Blossom	1 week
P2	Tauri desktop app PoC (Next.js static export + Tauri shell)	MO	2 weeks
P2	React Native shell wrapping current web views	D3	2 weeks

Section 9 — Domain Strategy (cubiqo.ai vs cubiqo.com) + Offline Browser Vision

9.1 cubiqo.ai vs cubiqo.com — The Definitive Recommendation

Short answer: Own both. Primary on .com . Brand on .ai .

Here is the full analysis:

Arguments for .com as primary:

- .com has ~48% of all registered domains — the default expectation globally
- Users type .com by muscle memory; .ai requires deliberate thought
- Trust signals: enterprise buyers, investors, and press default to .com
- Better email deliverability reputation (@cubiqo.com > @cubiqo.ai)
- SEO: Google treats .com with no special advantage, but legacy domain authority tends to concentrate there

- Insurance: if `.ai` ever becomes restricted (Anguilla controls it), `.com` is safe

Arguments for `.ai` as primary:

- Signal: `.ai` immediately communicates what the product is
- Memorability in the AI space: Perplexity.ai, Character.ai, Poe.ai all use it
- Differentiation from generic `.com` SaaS noise
- Premium positioning: `.ai` domains sell for 3-10x `.com` at resale
- Easier to get: cubiqo.com may be harder to secure if already taken

Recommended configuration:

```
Canonical domain: cubiqo.com      (all SEO juice lives here)
AI brand domain:  cubiqo.ai       (redirects 301 to cubiqo.com)
Alt spellings:    coqo.com / ciqo.com (redirect to cubiqo.com)
Email:           @cubiqo.com      (primary) + @cubiqo.ai (alias)
```

One exception: If cubiqo.com is NOT owned and cubiqo.ai IS — use `.ai` as primary immediately and try to acquire `.com` over time.

The 30+ country TLD strategy (from previous analysis):

Tier	Domains	Action	Reason
Protect	<code>.com</code> <code>.ai</code> <code>.io</code> <code>.co</code>	Own & keep	Core brand assets
Redirect	<code>.ca</code> <code>.uk</code> <code>.au</code> <code>.de</code> <code>.fr</code> <code>.in</code> <code>.sg</code> <code>.ae</code>	301 → <code>.com</code>	Key markets, affordable
Hold	<code>.sa</code> <code>.br</code> <code>.mx</code> <code>.za</code> <code>.ng</code> <code>.pk</code> <code>.ph</code> <code>.id</code>	Register if <\$50/yr	Emerging markets, option value
Release	All others	Let expire	No ROI, maintenance burden

9.2 The Offline Browser Vision — What Is Actually Feasible

The CEO described: an offline browser that uses AI search, works with user data offline, does photo/video editing, face tagging, AI model comparison, and data synthesis.

Let me split this into three categories: **Can do now**, **Can do in 6 months**, **Fundamentally hard or impossible**.

✅ CAN DO NOW (or within 8 weeks)

1. Offline-cached web pages (reading mode)

- Service worker already exists in `public/sw.js`
- Enhancement: add "Save for offline" button → caches full page HTML in IndexedDB
- User can read saved pages without internet
- Effort: 1 week

2. Local AI search over saved content

- Run keyword + semantic search over IndexedDB-cached pages
- Use existing TF-IDF or embed with a tiny local model (transformers.js, `all-MiniLM-L6-v2`, 80 MB)
- Effort: 2 weeks

3. Photo viewing + basic editing offline

- PWA can access files via File System Access API (Chrome 86+)
- Basic editing (crop, rotate, filter): use `fabric.js` or `jimp` (WASM port)
- No uploads needed — pure client-side
- Effort: 3 weeks

4. AI model answer comparison

- Already partially exists: `router.ts` routes between Claude, OpenAI, Ollama, MiniMax
- UI enhancement: "Compare" button sends same prompt to 2+ providers → side-by-side diff
- Offline: can compare WebLLM models (Phi-3 vs Gemma) without internet
- Effort: 1 week for UI, existing infrastructure handles routing

5. Local journal + memory synthesis

- Journal data is in Supabase BUT can be cached in `localStorage`/`IndexedDB` on write
- Offline journal: write to `IndexedDB` → sync to Supabase when back online
- Local synthesis: "Summarise my last 30 days" → WebLLM processes cached journal entries
- Effort: 2 weeks

⚠ FEASIBLE IN 6 MONTHS (with significant effort)

6. Face tagging in photos

- Technology: `face-api.js` (TensorFlow.js) — runs in browser, 97% accuracy on aligned faces
- Models: `ssd_mobilenetv1` (face detection) + `face_recognition_model` (128D embeddings)
- Workflow: scan photos → cluster by face vector similarity → user names clusters → persistent in `IndexedDB`
- Limitation: iOS Safari has poor WASM performance; desktop Chrome is excellent
- Effort: 4–6 weeks

7. Video editing (basic)

- `ffmpeg.wasm` runs in browser — can cut, trim, merge, add text overlays, change resolution
- File size limitation: WASM heap is ~2 GB max, so videos >500 MB need chunking
- Real-time preview is slow (not After Effects); good enough for social clips
- Effort: 3–4 weeks for basic editor UI

8. Proactive contextual recommendations (offline)

- User's data (journal, RGY capsules, saved pages) processed locally by WebLLM
- "Based on your journal this week, you seem to be focused on job hunting — want me to find openings?"
- Runs as a background Web Worker after WebLLM model is loaded
- Effort: 4 weeks

✗ FUNDAMENTALLY HARD OR NOT ADVISABLE

9. Offline web search (actual search engine)

- A real search index (like Google's) requires petabytes of data — impossible offline
- What IS possible: cached pages + local index of saved content (not "the web")
- Honest positioning: "Search your saved content offline" not "search the web offline"

10. Full video editing (After Effects level)

- GPU-accelerated video processing requires native code (Metal/CUDA) — not available in browser

- WASM is single-threaded-ish; real-time 4K editing is not feasible
- Tauri desktop app with FFmpeg native binary could handle this — separate effort

11. "Sending emails offline"

- SMTP requires internet — no workaround; emails queue locally and send when online
- This is fine and the right UX; just be honest about it in the UI

9.3 The "Cubiqo as AI Browser" Product Vision

Putting it together, here is what a 12-month roadmap for "Cubiqo as AI browser/OS" looks like:

PHASE 1 (Now → Launch): Core AI assistant + memory + PWA

- ✓ Chat, journal, agents, BYO keys
- ✓ PWA install on mobile/desktop
- ✓ Basic offline (cached pages, offline.html)

PHASE 2 (Month 1–3): Offline inference + saved content search

- WebLLM integration (Phi-3-mini in browser)
- "Save page" button → local AI search
- Offline journal with sync queue
- AI model comparison UI

PHASE 3 (Month 4–6): Media tools + face tagging

- Photo viewer + basic editing (crop/filter/text)
- Face clustering via face-api.js
- Video trimmer via ffmpeg.wasm
- Local media library (not cloud upload)

PHASE 4 (Month 7–12): Desktop + mobile native

- Tauri desktop app (macOS/Windows) + bundled Ollama
- React Native mobile app
- Full offline action queue (email/post queued, sent when online)
- Local file system access (journal exports, data backups)

Why this is the right direction — the WeChat analogy revisited:

WeChat succeeded not because it was a great chat app but because it became the infrastructure people lived in. The offline + local-first angle is Cubiqo's version of that:

- Cloud AI (Copilot, ChatGPT, Gemini): always online, always sending your data to servers
- Cubiqo offline mode: your AI runs on YOUR device, your data stays with YOU
- That is a genuine competitive moat — especially in privacy-conscious markets (EU, Canada, enterprise)

9.4 Implementation Priority Table

Feature	Feasibility	Effort	Revenue Impact	Privacy Story Value
Honest BYO UI copy	✓ Now	0.5 day	None	High (trust)
Offline action queue	✓ Now	1 week	Medium	High
Local journal cache	✓ Now	1 week	Medium	High
WebLLM integration	✓ 8 weeks	6 weeks	Very High	Very High

Feature	Feasibility	Effort	Revenue Impact	Privacy Story Value
AI model comparison UI	✅ 2 weeks	1 week	Medium	Medium
Photo editor (basic)	✅ 6 weeks	4 weeks	High	High
Face tagging	⚠️ 12 weeks	6 weeks	High	Very High
Video trimmer	⚠️ 12 weeks	4 weeks	Medium	Medium
Tauri desktop app	⚠️ 20 weeks	16 weeks	Very High	Very High
React Native app	⚠️ 24 weeks	20 weeks	Very High	Very High
"Search the web" offline	❌ Not viable	N/A	N/A	N/A

Section 10 — Who Are Our Users?

Grounded in actual Cubiqo feature set: Journal (Rozana), Job Hunt automation, RGY matching, Social Army, Emergent platform, in-app agents, BYO mode, Deals engine, Duo/Companion mode, CQ identity system, voice AI, crisis detection.

10.1 The Core Question — Why It Matters

"Who is the user?" is the most important question in product development. It determines:

- Which features to build next
- Which marketing channels to invest in
- What pricing model works
- What investor narrative to tell
- What the onboarding flow teaches first

Cubiqo is a **multi-persona product** — intentionally so, because it was built as an "AI OS" that different people use for entirely different reasons. That is a strength (broad TAM) and a risk (unfocused messaging). The answer is **to lead with one persona publicly, serve all privately**.

10.2 The Seven User Personas

PERSONA 1 — The Solopreneur Operator

"I run a one-person business and I need a full team inside an app."

Attribute	Detail
Age / Stage	28–45, already running a business or side hustle
Location	North America, UK, Australia, UAE, India (English-first markets)
Market size	64M solopreneurs in US alone; 400M+ globally (World Bank 2024)
Income	\$3K–\$25K/month revenue, under-resourced

Attribute	Detail
Tech comfort	Medium-high; uses Notion, Zapier, ChatGPT, Canva
Core pain	Wears every hat — marketing, ops, client work, admin — simultaneously
Cubiqo features used	Social Army (posts across 10 platforms), Agents (A1-A7 autonomous tasks), Deals engine (finds tools at discount), BYO mode (controls AI cost), Emergent (launches side projects), Journal (BigBoss accountability prompts)
Aha moment	Social Army auto-posts to 10 platforms while they sleep; morning shows 47 new impressions, zero effort
Monetization path	Free → Pro (\$29/mo) within 2 weeks once Social Army value clicks; upgrades to Founders Pass for unlimited agents
Acquisition channel	Twitter/X, Reddit r/entrepreneur, Indie Hackers, Product Hunt
Churn risk	Low — deeply embedded across social + ops workflows
Investor story value	★★★★★ — the "Zapier + Buffer + ChatGPT in one" narrative for the creator economy

PERSONA 2 — The Job Seeker in Transition

"I need a new job and the process is humiliating, manual, and exhausting."

Attribute	Detail
Age / Stage	24-40, recently laid off, career-pivoting, or quietly looking while employed
Location	US, Canada, UK, India, Australia
Market size	200M active job seekers globally at any moment (ILO 2024); US unemployment always $\geq$ 4M
Income	Variable — either unemployed or \$60K-\$120K employed and frustrated
Tech comfort	Medium; uses LinkedIn daily, spreadsheets to track applications
Core pain	Applies to 50+ jobs manually, hears nothing back, loses track, gets rejected for unknown reasons
Cubiqo features used	Job Hunt automation (resume upload → auto-apply → status tracking), LinkedIn integration, cover letter AI, interview coaching via voice AI, Journal for rejection processing
Aha moment	Uploads resume Monday, wakes up Tuesday to 12 applications submitted + 2 interview alerts in inbox
Monetization path	Free tier for 10 apps/month → Pro (\$29/mo) for unlimited auto-apply; crisis persona is emotionally motivated to pay
Acquisition channel	Reddit r/cscareerquestions r/jobs, LinkedIn job seeker groups, YouTube career coaches, TikTok "job hunt tips"
Churn risk	

Attribute	Detail
	Medium — leaves when they get a job (natural churn); re-acquires on next career move
Investor story value	★★★★★ — automated job applications is a hot category; ZipRecruiter is a \$2B company

Feature gap to fix: Full autonomous application submission (currently tracks but doesn't auto-submit). This single gap is the difference between "interesting tool" and "life-changing product" for this persona.

PERSONA 3 — The Emotionally Overwhelmed Professional

"I am doing okay on paper but I am exhausted, isolated, and have no one to actually talk to."

Attribute	Detail
Age / Stage	27–45, high-functioning but emotionally depleted
Location	Urban North America, UK, Australia — high-cost, high-pressure cities
Market size	76% of US workers report burnout (Deloitte 2023); 1 in 5 adults experience mental health challenge annually
Income	\$60K–\$200K — can afford a therapist but won't commit to one
Tech comfort	Medium; smartphone-heavy, app-native
Core pain	No safe space to process stress; therapy stigma or cost barrier; journaling feels juvenile alone
Cubiqo features used	Journal/Rozana (BigBoss "friend who tells you the truth" prompts), Duo mode (side-by-side companion), voice AI (emotionally modulated responses), crisis detection (self-harm escalation to safe resources), Companion mode
Aha moment	Journal prompt "Your boss just took credit for your work. Say the unsayable." — they type something raw and Cubiqo responds with exactly the right tone: validating, then redirecting
Monetization path	Hardest to convert on price alone → convert via habit (daily journal streak, weekly emotional summary email) → upgrade for voice companion + unlimited history
Acquisition channel	Instagram mental health accounts, Reddit r/mentalhealth r/therapy, therapy-adjacent podcasts, TikTok therapist creators
Churn risk	Very low once habit forms (daily journal creates daily return)
Investor story value	★★★★★ — mental health tech is a \$6B market but requires careful positioning to avoid regulatory risk; better to position as "emotional productivity" than "therapy"

Critical warning: This persona triggers real crisis situations. The crisis detection escalation path (policy-router.ts) MUST be battle-tested before this persona is publicly targeted. Mental health claims in marketing require legal review.

PERSONA 4 — The Privacy-Conscious Tech Enthusiast

"I want powerful AI but I refuse to send my data to Big Tech servers."

Attribute	Detail
Age / Stage	25–45, developer or highly technical non-developer
Location	EU (GDPR-native), Canada, Germany, Switzerland, Japan
Market size	Smaller absolute numbers (~5M globally) but extremely vocal and viral
Income	\$80K–\$200K, willing to pay premium for privacy
Tech comfort	High — runs their own servers, uses Signal, has a VPN, knows what RAG means
Core pain	Every AI product sends their data to OpenAI/Google; no true privacy option exists at a consumer price point
Cubiqo features used	BYO mode (own API key, controls data routing), future WebLLM (runs inference entirely in browser), Ollama integration, privacy-first journal
Aha moment	Sees BYO settings page — "This is the first AI app that lets me use MY key, MY model, and tells me exactly what happens to my data"
Monetization path	Premium/Enterprise tier — this persona pays for the platform fee + will pay for Founders Pass as status symbol
Acquisition channel	Hacker News, r/privacy, r/selfhosted, Mastodon, DEF CON community, Lobsters
Churn risk	Very low — will become an evangelist and pull in 3–5 more users through word of mouth
Investor story value	★★★★★ — EU AI Act compliance + privacy-first positioning is increasingly a competitive moat; powerful for EU expansion narrative

Note: Current BYO still routes through Cubiqo server (only key is user's own). True privacy requires WebLLM (Tier 2, 8 weeks). Do NOT market to this persona as "private" until Tier 2 ships — they will call it out publicly and cause reputational damage.

PERSONA 5 — The Creator / Social Media Entrepreneur

"My content is my product and I can't keep up with the pace of posting."

Attribute	Detail
Age / Stage	20–38, full-time or aspiring creator / influencer
Location	US, UK, Brazil, India, Philippines, Nigeria
Market size	50M+ content creators globally (SignalFire 2023); 2M+ earn full-time income from content
Income	\$1K–\$30K/month from content; time-constrained and inconsistent
Tech comfort	High for social tools; medium for dev tools

Attribute	Detail
Core pain	Content creation is a daily grind; one person cannot maintain 5+ platforms without burning out
Cubiqo features used	Social Army (10 platforms × 10 accounts × scheduled posting), GFXToolz integration (images/videos), AI caption generation, persona management (different voice per platform), Deals (sponsored content deal tracking)
Aha moment	Sets up Social Army once, schedules a week of posts across Instagram/TikTok/ Twitter/LinkedIn/YouTube — all from one dashboard
Monetization path	Pro immediately (\$29/mo is trivial if it saves 10 hours/week); potential Founders Pass for white-label client management
Acquisition channel	YouTube creator economy content, CreatorIQ, Creator Commerce, Beehiiv newsletters for creators
Churn risk	Medium — churns if Social Army hits API platform restrictions (LinkedIn 1/day, Instagram approval delays)
Investor story value	★★★★★ — creator economy is the hottest B2C market; Buffer is worth \$80M, Hootsuite \$750M; Cubiqo is Buffer + AI agents

PERSONA 6 — The Developer / Technical Builder

"I want to build and ship products faster without managing 20 different tools."

Attribute	Detail
Age / Stage	22–40, freelance dev, indie hacker, startup engineer
Location	Global (English-first); strong in India, Eastern Europe, Southeast Asia, Canada
Market size	28M developers globally (Evans Data 2024); 4M freelance devs in US
Income	\$40K–\$200K; cost-sensitive on tools, but pays for 10x productivity gains
Tech comfort	Very high — daily VS Code, GitHub, terminal users
Core pain	Too many disconnected tools: IDE + hosting + secrets management + deployment + monitoring = cognitive overhead
Cubiqo features used	Emergent platform (Monaco editor + sandbox + deploy pipeline), BYO mode (own OpenRouter/Groq key for cheap inference), Agents (code review, test generation, deployment), CQ social for dev community
Aha moment	Opens Emergent, types a project description, watches Cubiqo scaffold a Next.js app, configure secrets, and deploy to Vercel — in one session
Monetization path	Heavy free user first (builds on Emergent); converts to Pro when projects earn money or team joins; strong Founders Pass candidate for agency use
Acquisition channel	Hacker News "Show HN", GitHub trending, DEV.to, r/programming, YouTube dev channels (Theo, Fireship)
Churn risk	Low once Emergent hosts a live project (switching cost = redeployment effort)

Attribute	Detail
Investor story value	★★★★★ — Vercel, Railway, Replit all multi-billion; Cubiqo is the AI-native version of all three

PERSONA 7 — The RGY Community Seeker

"I want to connect with people who are actually on a similar wavelength, not just network for networking's sake."

Attribute	Detail
Age / Stage	22-45, socially active but disillusioned with LinkedIn/Twitter networking
Location	Urban global; early adopter in tech / creative / entrepreneurship communities
Market size	Hard to size directly — this is the "quality connection" persona; target is 50M LinkedIn-frustrated professionals
Income	Variable — defined more by mindset than income
Tech comfort	Medium-high
Core pain	Networking feels transactional; LinkedIn is a spam machine; can't find genuine collaborators or people in the same life chapter
Cubiqo features used	RGY chat (express intent in colour-coded context: Red=transform, Yellow=grow, Green=stabilise), capsule matching (intent × keyword × vector similarity → anonymous match → consent → reveal), CQ number as identity, CQ-to-CQ direct messaging with voice delivery
Aha moment	Sets RGY intent to Yellow: "looking for a technical co-founder for a fintech idea in Toronto" → Cubiqo surfaces 3 anonymised matches → all 3 are genuinely relevant → accepts 1 → that match becomes a co-founder
Monetization path	Free RGY → Pro for unlimited capsules + proactive discovery + video RGY rooms → Premium for concierge matching
Acquisition channel	RGY is the viral hook — shareable CQ number, RGY colour profile share cards, community subreddits, co-founder matching communities (YC co-founder matching, Indie Hackers)
Churn risk	Very low — active matching creates weekly return visit; matched connections create emotional ownership of the platform
Investor story value	★★★★★★ — LinkedIn was a \$26B acquisition; the "next LinkedIn" narrative plus patentable matching algorithm (Patent 3) makes this the single most investor-compelling feature

10.3 Persona × Feature Matrix

Feature	Solopreneur	Job Seeker	Emotional	Privacy	Creator	Developer	RGY Seeker
Journal / Rozana	✓	✓	★	✓	—	—	✓
Job Hunt Automation	—	★	—	—	—	✓	—
Social Army	★	—	—	—	★	—	✓
RGY Matching	✓	✓	✓	—	✓	✓	★
In-App Agents	★	✓	—	✓	✓	★	—
BYO Mode	✓	—	—	★	—	★	—
Emergent Platform	✓	—	—	✓	—	★	—
Voice AI	✓	✓	★	—	✓	—	✓
Deals Engine	★	—	—	—	✓	✓	—
Duo / Companion	✓	✓	★	—	—	—	✓
CQ Number / Identity	✓	✓	✓	✓	✓	✓	★
Crisis Detection	—	—	★	—	—	—	—

★ = primary feature for this persona | ✓ = secondary use

10.4 Which Persona to Lead With at Launch

Lead publicly with: Persona 1 (Solopreneur Operator) + Persona 5 (Creator)

Reason: These two personas have:

- Highest willingness-to-pay (already spending on Buffer, Zapier, ChatGPT)
- Lowest regulatory risk (no mental health or employment law exposure)
- Most viral distribution (they POST publicly, so Cubiqo gets organic impressions)
- Clearest ROI story ("saves 10 hours/week, worth \$29/mo")
- Strongest investor narrative (creator economy + solopreneur market)

Serve privately at launch, market loudly at Month 3:

- Persona 2 (Job Seeker) — once auto-apply is complete
- Persona 7 (RGY Community) — once capsule matching has 100+ users to match against

Serve carefully, never lead marketing with:

- Persona 3 (Emotional) — position as "emotional productivity" not "mental health tool"; needs legal review first
- Persona 4 (Privacy) — position as "data control" not "privacy"; needs WebLLM shipped first

10.5 The ONE User Who Represents All Seven

If forced to pick a single archetype for everything — product, marketing, pitch:

"Alex, 32, runs a 2-person design + copywriting agency in Toronto. Alex uses Cubiqo to post across 8 social platforms daily, track 3 client projects inside Emergent, decompress after client calls in the journal, and has an RGY Yellow capsule active looking for a dev co-founder to productize the agency's workflows. Alex pays \$29/month and would pay \$99/month if the voice companion was smarter."

Alex is real. Alex is everywhere. Alex is the user Cubiqo was built for.

10.6 What This Means for Onboarding

The onboarding flow MUST ask one question first:

"What matters most to you right now?"

- | | |
|---------------------------------------|----------------------------------|
| A) Growing my business / brand online | → routes to Social Army + Agents |
| B) Finding a new job | → routes to Job Hunt setup |
| C) Building something new | → routes to Emergent + Agents |
| D) Making sense of my day / life | → routes to Journal (Rozana) |
| E) Connecting with the right people | → routes to RGY intent setup |

Each answer triggers a different onboarding sequence, a different first Aha moment, and a different first 30-day email series. This single onboarding branching decision will likely double activation rate versus a one-size-fits-all tour.

Current gap: Onboarding is currently a single generic flow. This branching is a P1 priority — build it in week 3–4 post-launch.

10.7 Market Size Summary

Persona	Addressable Market	Realistic TAM (Cubiqo serves)	Conversion at 1%
Solopreneur Operator	400M globally	20M English-first, tech-ready	200K users
Job Seeker	200M active seekers	10M English-first, tech-ready	100K users
Emotional Professional	500M+ burnout-affected	15M willing to use AI for this	150K users
Privacy Enthusiast	30M privacy-conscious	5M willing to pay premium	50K users

Persona	Addressable Market	Realistic TAM (Cubiqo serves)	Conversion at 1%
Creator / Social	50M creators	5M with posting automation need	50K users
Developer / Builder	28M developers	8M indie/freelance	80K users
RGY Community Seeker	900M LinkedIn users	10M disillusioned networkers	100K users

Combined realistic target at 1% TAM conversion: 730,000 users

At \$29/month average blended ARPU: \$21.2M ARR

At \$15/month average blended ARPU (freemium mix): \$131M ARR potential at full TAM

This is a \$100M+ ARR opportunity if Cubiqo executes across all seven personas over 5 years. The near-term focus should be 1,000 users → 10,000 users → first \$100K ARR. All three are achievable within 18 months with the right execution.

Section 11 — Flagship Spec Gap Analysis: Where Does Cubiqo Stand?

Reference document: The "CUBIQO — Flagship" product specification provided February 2026.

Method: Every claim below is grounded in direct code inspection of the branch `copilot/investigate-features-and-ui-components`. No assumptions.

Legend: ● Complete · ● Partial (skeleton/type/UI only, core not fully wired) · ● Missing · ⚠ Deviation (works but differently than spec)

11.1 Overall Scorecard

Spec Section	Status	Short Verdict
Color / Voice (UI-level)	●	TEAL/RED/YELLOW defined with voice tones, lock commands, and per-color animation parameters
Input / Output	●	STT ☒ , TTS ☒ , state machine ☒ ; Vocspad input mode ☒ not found anywhere
State Machine (L→T→S→I)	●	<code>useAIState.ts</code> implements all 4 states with event logging
UI/UX — Wireframe cuboid	●	Multiple variants: <code>EnergyCubeWireframe</code> , <code>SilverWireLandingCube</code> , <code>IsometricCube</code>
UI/UX — Solid cuboid	●	<code>AICuboidGLB.tsx</code> (GLB-loaded, InnerPlasma + OuterGlass), <code>SettingsCube.tsx</code> (RoundedBox)
	●	GLB has OuterGlass + color; ≤50% color rule NOT enforced

Spec Section	Status	Short Verdict
UI/UX — Hybrid material + color		
UI/UX — Timing tokens	●	<code>--duration-normal: 200ms</code> in CSS; glow 150-300ms tokens not formalized
UI/UX — Reduced motion	●	<code>motion-safe</code> : Tailwind used in VoiceStateIndicator; not applied to 3D cube transitions
Special Moves (8 total)	●	5 of 8 animated (Resonance, Breakthrough, Memory Thread, Deep Focus, Co-Presence); Wink, Trust Earned, Handoff are missing
Audio Cues (wake/tick/DND)	●	AudioContext manager exists for playback; zero wake chime, zero speak ticks, no DND mode
Side Panel — Keywords	●	<code>KeywordPanel.tsx</code> with per-color RGY lists, tap-to-edit, session-scoped localStorage
Side Panel — Geo-fence	●	Referenced in <code>capsule-manager.ts</code> and <code>discovery-service.ts</code> ; not surfaced in UI
Side Panel — Color lock	●	<code>cubiqo.color.lock('RED')</code> command in <code>SettingsCube/commands.ts</code>
SettingsCube — Voice config	●	<code>cubiqo.voice.set()</code> command; ElevenLabs per-color voice tones defined
SettingsCube — Live confirm	●	CommandInput shows result; no spoken live confirmation
SettingsCube — Event logs	●	Console logging only; no structured event-only log store
RGY Router — Backends	●	GPT-4o (OpenAI), Claude 3.5 (Anthropic), Ollama (self-hosted), OpenRouter aggregation
RGY Router — Intent routing	❖❖	Zone detection → model selection in <code>policy-router.ts</code>
RGY Router — Auto-failover	●	3-level failover per zone (primary → secondary → ultimate fallback)
RGY Router — Scoped to CUBIQO	●	No <code>/rgy/route</code> endpoint; routing is embedded in chat API; Worlds boundary undefined
RGY Router — CAP Orchestrator	●	Not implemented; no file, no route, no reference to this concept
RGY Router — Keyword feedback	●	Session-scoped localStorage only; no DB telemetry pipeline back from sub-domains
RGY Router — Direct override	●	<code>isFounder</code> flag routes to Sonnet; no explicit <code>model:gpt claude local</code> UI toggle
RGY Router — Zero retention	⚠	"Zero-Retention" is marketing copy in email template; memory system STORES to Supabase; contradiction
Feature — BYO Mode	●	

Spec Section	Status	Short Verdict
		byo-manager.ts , BY0Settings.tsx , /api/byo/route.ts , AES-256 key encryption
Feature — Auth (magic link)	●	Supabase magic link, /auth/callback
Feature — Auth (WebAuthn/Passkeys)	●	4 WebAuthn routes, BiometricRegistration.tsx , BiometricLogin.tsx
Feature — Email	●	Resend integration, email templates, /api/channels/[type]
Feature — Calendar	●	Channel API exists; no calendar-specific implementation found
Feature — Food delivery	●	action-types.ts defines ubereats/doordash platform types; no live API calls
Feature — Taxi / ride-share	●	Mentioned in action-types; no live integration
Feature — Browser automation	●	Puppeteer-based BrowserService.ts with navigate, click, type, fill, screenshot, extract
Feature — Smart-home	●	integration-registry.ts lists smart_home type; no implementation (Google Home, HomeKit, MQTT)
Feature — Intelligent chat/match	●	RGY capsule system, discovery cron, proactive matching service
Feature — CQ↔CQ connections	●	cq-to-cq/ module: CQ number generator, WebSocket, WebRTC calls — not connected to UI
Feature — Wallet / QR escrow	●	wallet-service.ts with QR delayed release backend; no Stripe, no UI, no DB migration
RGY Sub-domain — No coupling	●	Keywords session-scoped; sub-domains do not force route changes
RGY Sub-domain — CAP Orchestrator	●	As above — not built
RGY Sub-domain — Telemetry hint	●	localStorage session hint; not a structured telemetry pipeline

11.2 Detailed Gap Analysis by Spec Section

COLOR / VOICE (UI-LEVEL) — ● Substantially Complete

What the spec requires:

- TEAL (goal-oriented, professional/decisive voice)
- RED (age-gated/explicit, discreet/low-volume voice)
- YELLOW (casual/general, friendly, light sarcasm)
- Color/voice are operational signals; do NOT select models or imply emotion

- Self-harm: force YELLOW, block instructions, offer resources
- User may lock to one color + voice; no cross voice↔color mixing

What exists in code:

src/config/colors.ts defines all four colors (RED/YELLOW/TEAL/ORANGE) with per-color:

- voiceTone strings (e.g. TEAL: "clear, motivating, balanced"; RED: "deep, slow, sensual whisper; discreet")
- animationSpeed , glowIntensity , breathingSpeed , blinkStyle

src/lib/ai/policy-router.ts correctly forces YELLOW zone and adds support prompt on self-harm detection:

```
if (selfHarmPatterns.test(lastMessageText)) {
  zone = 'YELLOW';
  systemPrompt += " IMPORTANT: ... extremely supportive ... provide help resources.";
}
```

src/lib/settings-cube/commands.ts supports:

```
case 'cubiqo.color.lock': { const color = ... if (['RED','YELLOW','TEAL','ORANGE'].includes(color)) ... }
```

Gaps:

1. **Color as model selector:** Policy-router DOES route RED zone to uncensored models (MIXTRAL_8X22B , LLAMA_UNCENSORED). This contradicts the spec rule: "Color does not select models." RED is the policy zone for age-gated content — but the spec says color is UI-only and routing is purely intent-based. **Spec vs implementation mismatch.** ⚠️
2. **Age-gate / age verification:** colors.ts says emotion: 'age-gated, critical, goal-oriented' for RED, but there is **zero age verification code** anywhere. The RED zone is accessible to all users with no gate.
3. **ORANGE color:** Not in the spec (spec has TEAL/RED/YELLOW only). ORANGE is a Fourth-Way philosophical color that Cubiqo added. This is an extension, not a violation.

Fix required:

- Add age-gate middleware check before allowing RED zone access (simple: DOB confirmation modal stored to user profile)
- Document that RED zone uses specific models as a deliberate design decision (or refactor router to be model-agnostic per-zone, using only intent for model scoring)

INPUT / OUTPUT — 🟡 Partial

What the spec requires:

- Input: speech OR **Vocspad** (type or talk)
- Output: on-screen text + synthesized speech; synchronized with status cues
- State machine: Listening → Thinking → Speaking → Idle

What exists:

Requirement	Code Evidence	Status
Speech input (STT)	/api/stt/route.ts using Whisper	🟢

Requirement	Code Evidence	Status
Text input (chat)	ChatContainer.tsx , /api/chat/route.ts	●
Vocspad (combined type+talk mode)	✗ No file named Vocspad anywhere in 300+ files	●
TTS output	/api/tts/route.ts , ElevenLabs, OpenAI TTS fallback	●
On-screen text	Chat bubbles in ChatContainer.tsx	●
Synchronized status cues	VoiceStateIndicator.tsx with pulse rings	●
State machine	useAIState.ts (idle/listening/thinking/speaking)	●

Vocspad gap: The spec treats Vocspad as a specific named input mode — a hybrid voice+text pad that lets users type and talk simultaneously. This UX component does not exist. The current chat has separate text input and a mic button, but not a unified Vocspad component that merges both modes in a single surface.

Fix required (2 weeks):

Create `src/components/chat/Vocspad.tsx` — a unified input component that:

- Shows a text area with live STT transcription overlaid
- Has a mic toggle button embedded inline
- Falls back gracefully to text-only when no mic permission

UI/UX — CUBOID FORMS & MATERIALS — ● Partial

What the spec requires:

Form	Description	Code File	Status
Outline-only cuboid (fig1)	Floating isometric wireframe, crisp edges, no fill	EnergyCubeWireframe.tsx , SilverWireLandingCube.tsx	●
Solid cuboid (fig2)	Filled body, subtle depth, no interior facets	AICuboidGLB.tsx (OuterGlass solid), SettingsCube.tsx (RoundedBox)	●
Hybrid material + color (fig3)	Material (e.g. metal/glass) + operational color overlay	AICuboidGLB.tsx (OuterGlass + InnerPlasma + color param)	●
≤50% color visible area rule	Hero/transition hybrid must show ≤50% color	✗ Not enforced	●
Glass material	For aesthetics only, high contrast	GlassyAgentCube.tsx , NeonGlassCube.tsx	●
Metal material	For aesthetics only	SilverWireLandingCube.tsx (chrome/silver)	●
Fabric-soft-touch material	For aesthetics only	✗ Not implemented	●

Form	Description	Code File	Status
Wireframe/Neon Outline variant	Compact/overlays, UI-only, router unchanged	EnergyCubeWireframe.tsx , neon demo page	●
Split Material/ Color hero variant	≤50% color coverage	AICuboidGLB partial	●

Timing Tokens:

Token	Spec	Code	Status
Color/material swap	≤200ms	<code>--duration-normal: 200ms</code> in <code>globals.css</code>	●
Glow-in/glow-out	150–300ms	Not formalized as CSS variable	●
Reduced motion respect	All animations	<code>motion-safe:</code> in <code>VoiceStateIndicator</code> ; missing on 3D cube lerp transitions	●

Fabric-soft-touch fix (1 week): Create a `FabricCube.tsx` variant using a `Three.js MeshStandardMaterial` with low metalness, high roughness, and a subtle normal map to simulate soft textile.

≤50% color rule enforcement: Add a CSS custom property `--cube-color-coverage` and enforce via a shader uniform `uColorCoverage` clamped to 0.5 in the GLB material.

SPECIAL MOVES — ● 5 of 8 Implemented

What the spec requires: 8 named Special Moves, all UI-only, respecting timing tokens.

Move	Spec	Code in <code>AICuboidGLB.tsx</code>	Status
Resonance	Rapid pulse	<code>emissiveModifier = 1.0 + Math.sin(t * 10) * 0.5</code>	●
Breakthrough	Flash	<code>(Math.floor(t * 20) % 2) ? 1.5 : 0.5</code>	●
Trust Earned	(unspecified visual)	✗ Not in switch statement	●
Co-Presence	(unspecified visual)	<code>innerRef.current.rotation.z = t * 2</code>	●
Wink	YELLOW-only	✗ Not in switch statement	●
Deep Focus	(unspecified visual)	<code>scale.setScalar(0.8 + Math.sin(t * 2) * 0.1)</code>	●
Memory Thread	(unspecified visual)	<code>rotationMultiplier = 5.0</code>	●
Handoff	(unspecified visual)	✗ Not in switch statement	●

Fix required (1 week): Add 3 missing cases to `AICuboidGLB.tsx` animation switch:

```
case 'Trust Earned':
  // Warm golden pulse + scale breathe – acknowledgement moment
  mat.emissive.lerp(new THREE.Color('#ffd700'), 0.1)
  innerRef.current.scale.setScalar(1.0 + Math.sin(t * 3) * 0.05)
  break
case 'Wink':
  // YELLOW-only: quick rotation tilt + emissive blink, 200ms duration
  if (color === 'YELLOW') {
    groupRef.current?.rotation.set(0, 0, Math.sin(t * 30) * 0.15)
    emissiveModifier = Math.sin(t * 40) > 0 ? 1.3 : 0.7
  }
  break
case 'Handoff':
  // Outward scale + fade → signals transition to another agent/world
  innerRef.current.scale.setScalar(1.0 + moveTimerRef.current * 0.3)
  mat.opacity = Math.max(0, 1.0 - moveTimerRef.current * 0.5)
  break
```

AUDIO CUES — ● Not Implemented

What the spec requires:

- **Wake:** brief chime on activation (optional haptics)
- **Speak start/stop:** soft ticks aligned to TTS start and end
- **Error/Alert:** single neutral tick
- **Controls:** volume slider, on/off toggle, DND mode

What exists:

- `src/lib/audio/audioContext.ts`: AudioContext manager — handles playback unlocking on user gesture. No sound files, no generated tones.
- `src/lib/audio/audio-score-service.ts`: Background ambient scoring — oscillators for ambient music. No UI cues.
- `src/lib/multimodal/audio.ts`: STT recording helpers. No cue sounds.

Nothing implements any of the required audio cues.

Fix required (2 weeks):

Create `src/lib/audio/ui-cues.ts`:

```
// Uses Web Audio API oscillators – zero external files needed
export function playWakeChime() { /* 220Hz → 440Hz, 80ms, ramp out */ }
export function playSpeakTick() { /* 880Hz, 30ms, soft envelope */ }
export function playErrorTick() { /* 330Hz, 50ms, neutral */ }

// DND mode – stored to user preferences
export function setDNDAudio(enabled: boolean) { ... }
export function setAudioVolume(level: 0 | 0.25 | 0.5 | 0.75 | 1.0) { ... }
```

Call sites:

- `useAIState.ts`: call `playWakeChime()` on `idle` → `listening` transition
- `useElevenLabsTTS.ts`: call `playSpeakTick()` on TTS stream start/end
- Chat error handler: call `playErrorTick()` on AI failure

SIDE PANEL & SETTINGSCUBE — Substantially Complete (2 minor gaps)

What the spec requires + code evidence:

Requirement	Code	Status
Keywords panel with per-color lists	<code>KeywordPanel.tsx</code> — RGY cards with tap-to-edit and session keywords	
User can add intent keywords	Add/remove in <code>KeywordPanel</code> ; stored to <code>localStorage</code> per session	
Company / collaboration / trade intents	<code>CARD_CONFIG</code> in <code>RGYColorSelector.tsx</code> maps contexts	
Geo-fence for match	<code>capsule-manager.ts</code> and <code>discovery-service.ts</code> reference geofence	 (not in UI)
Color lock	<code>cubiqo.color.lock('TEAL')</code> in <code>SettingsCube</code> commands	
Voice lock with color	<code>cubiqo.voice.set({...})</code> command	
No cross voice↔color mixing	Lock command sets both; no mixing API exposed	
Speak to update config	<code>SettingsCubeApp</code> accepts typed/voiced commands	
Live confirmations	<code>CommandInput</code> shows result text; no spoken TTS confirmation	
Event-only logs	<code>Console.log</code> only; no structured event store	

Gap: Geo-fence UI. The geofence data model exists in `rgy_capsules` table (from previous migration analysis) but the `KeywordPanel` and RGY settings do not expose a "set my location radius" control. Fix: add a geo-fence radius slider to `ProMatchSettings.tsx`.

Gap: Spoken live confirmation. When a user says "lock to TEAL," the cube changes color but `Cubiqo` does not speak a confirmation back. Fix: after `executeCommand()` succeeds in `SettingsCubeApp.tsx`, call the TTS API with the confirmation text.

RGY ROUTER — Mostly Complete; CAP Orchestrator and Zero-Retention Critical Gaps

Backends and failover:

Requirement	Code	Status
GPT backend	<code>callOpenAI()</code> in <code>router.ts</code> , <code>OpenRouter</code> GPT-4o in <code>policy-router.ts</code>	
Claude backend	<code>callClaude()</code> in <code>router.ts</code> , <code>Claude 3.5 Sonnet</code> in <code>policy-router.ts</code>	
Self-hosted LLM (Ollama)	<code>callOllamaWithFallback()</code> , <code>isOllamaAvailable()</code> in <code>ollama.ts</code>	
Auto-failover on errors	3-level try/catch chains per zone in <code>policy-router.ts</code>	

Requirement	Code	Status
SLA breach failover	✗ No SLA timer / timeout threshold	●
Direct override (model:gpt\ claude\ local)	isFounder → Sonnet; no explicit per-message model: field in UI	●
Reasoning flag	reasoning: true → DeepSeek R1 → Opus	●

CAP Orchestrator — ● Critical Gap:

The spec defines a "CAP Orchestrator" that Worlds (sub-domains/tools) use, inheriting the chosen backend without calling `/rgy/route` directly. This concept does not exist in the codebase. There is:

- No `cap-orchestrator.ts` file
- No `/api/cap/` routes
- No "World" abstraction in code
- The agents system (`engine/agent.ts`) is the closest concept but it routes to external LLMs directly, bypassing zone/policy logic entirely

This is the **largest architectural gap** in the spec. The CAP Orchestrator is the integration boundary that would allow Emergent, Social Army, Job Hunt, Journal and other sub-domains to route through the same policy-aware backbone as the main Cubiqo chat. Without it, every sub-domain calls its own LLM independently with no shared zone policy, no shared keyword telemetry pipeline, and no shared safety guardrails.

Build plan for CAP Orchestrator (4 weeks):

```

Week 1: Define `CAPOrchestrator` class in `src/lib/cap/orchestrator.ts`
  - Accepts: worldId, userId, messages, zoneHint?
  - Inherits active zone from user session
  - Passes through PolicyRouter with CUBIQO-scoped config
Week 2: Expose `/api/cap/route` endpoint
  - All internal tools/Worlds call this instead of OpenAI directly
Week 3: Migrate Emergent, Social Army to use CAP endpoint
Week 4: Add keyword telemetry pipeline: CAP → session store → Side Panel push

```

Zero-Retention Contradiction — ⚠ Spec vs Implementation:

The spec states: "Zero retention; contextual recommendations only; adaptive within session; no stored profiles."

But the codebase has:

- `/api/memory/route.ts` — stores and retrieves memory from Supabase
- `/api/extract-memories/route.ts` — extracts and persists memories from conversations
- `/api/journey/memories/route.ts` — journey memory with embeddings stored in DB
- `magic-link.ts` email template says "Zero-Retention. Private." as marketing

This is a fundamental contradiction. The product's memory system IS its competitive differentiator (conscious memory, journey tracking). The spec's "zero retention" language was likely written for the ROUTING layer (router doesn't store conversation content), not the entire product.

Resolution needed: Clarify the spec to say "Zero retention at the routing layer" — the router itself does not store conversation data. The conscious memory system, when enabled by the user, stores memories with explicit consent. Update the privacy policy accordingly.

FEATURE SET — Status by Feature

Feature	Spec Requirement	Code Evidence	Status	Gap
BYO Mode	Cloud + API keys	<code>byo-manager.ts</code> , <code>/api/byo/</code> , AES-256	●	None
Magic-link auth	Passwordless	Supabase + Resend templates	●	None
OAuth / OIDC	Account creation	Supabase OAuth providers	●	None
Passkeys	WebAuthn	4 WebAuthn routes + UI components	●	None
Email	Send/receive	Resend + channel API	●	Receive via webhook only
Calendar	Read/write events	Channel type exists; no iCal/Google Calendar API	●	3–4 weeks to wire Google/Outlook Calendar API
Food delivery	Order placement	<code>action-types.ts</code> defines platform types; <code>HandshakeWizard.tsx</code> shows Uber Eats	●	No live API calls; needs UberEats/DoorDash developer account
Taxi / ride-share	Book rides	<code>action-types.ts</code> type only	●	Uber/Lyft API or browser automation fallback
Browser access	Navigate/interact	<code>BrowserService.ts</code> Puppeteer, <code>/api/browser/</code>	●	Puppeteer not available on Vercel Edge; needs server
Smart-home control	Device control	Integration registry type only	●	Needs Google Home/HomeKit/MQTT SDK
Browser automation (tickets)	Form fill, navigate	<code>BrowserService</code> <code>navigate+click+fillForm</code>	●	Consent gate wired correctly
Intelligent chat/match	RGY capsule matching	Capsule manager + discovery cron + RGY rooms	●	Proactive matching not user-triggered
CQ↔CQ permanent connections	P2P messaging	<code>cq-to-cq/</code> module: CQ number, WebSocket, WebRTC	●	Module exists; not wired to main chat UI
Geo-fence for match	Location-aware matching	Referenced in capsule-manager; no UI control	●	Geo-fence radius picker missing from Side Panel

Feature	Spec Requirement	Code Evidence	Status	Gap
Wallet payments	Crypto + fiat	wallet-service.ts with escrow	●	No Stripe; no crypto payment gateway (MetaMask/ Coinbase)
QR-based delayed release	Escrow via QR	wallet-service.ts — QR code + held/released status	●	DB migration for payments table not found

11.3 Top Priority Gaps to Close Before Launch

Priority	Gap	Effort	Risk if Skipped
● P0	Age gate for RED zone access	1 week	Legal liability if minors access explicit content
● P0	Zero-retention spec clarification + privacy policy update	3 days	Regulatory (GDPR) and trust damage if users believe no data is stored but it is
● P0	Smart-home: spec lists it as a feature	3 weeks	False advertising if listed in feature set at launch
● P1	CAP Orchestrator (routing unification)	4 weeks	Sub-domains bypass safety guardrails; keyword telemetry broken
● P1	Audio cues (wake, tick, DND)	2 weeks	Poor UX; voice mode feels unpolished
● P1	Vocspad unified input	2 weeks	Core UX element missing from spec
● P1	Special Moves 3 missing (Wink, Trust Earned, Handoff)	1 week	Incomplete product spec delivery
● P1	CQ↔CQ connection to main chat UI	2 weeks	CQ number system exists but users can't use it
● P2	Wallet DB migration + Stripe	3 weeks	No monetization beyond CubiKey
● P2	Calendar API (Google/Outlook)	3 weeks	Email without calendar is half a productivity suite
● P2	Spoken SettingsCube confirmations	3 days	Minor UX polish
● P2	Geo-fence UI in Side Panel	1 week	RGY match is less accurate without user location

11.4 What Is Actually Production-Ready Against This Spec

Despite the gaps, the following flagship-spec components are **genuinely production-ready**:

- 1. **Policy Router** — zone detection, 3-backend failover, self-harm safety, language adaptation, founder mode, reasoning path, freedom path, search-all. This is solid.
- 2. **Color system** — TEAL/RED/YELLOW with per-color voice tones, animation parameters, lock commands. Fully wired.
- 3. **State machine** — 4-state (idle/listening/thinking/speaking) with event logging.
- 4. **Cuboid 3D forms** — wireframe and solid variants working; GLB-loaded AICuboidGLB with 5 Special Moves.
- 5. **STT + TTS pipeline** — Whisper STT, ElevenLabs TTS with voice modulation (Madhyama Marg), OpenAI TTS fallback.
- 6. **BYO Mode** — complete with AES-256 key encryption, per-model routing, test endpoint.
- 7. **Auth** — magic link, OAuth/OIDC, WebAuthn/Passkeys — all four auth modes.
- 8. **Browser automation** — Puppeteer service with consent gate, form fill, navigate, screenshot.
- 9. **RGY capsule matching** — 4-signal matching (colour → intent → keyword → vector), proactive discovery cron.
- 10. **Side Panel keywords** — per-color RGY keyword lists with session persistence.

11.5 Spec Alignment Summary

SPEC COVERAGE: ~63% Complete

Completely built:	12 / 33 requirements		36%
Partially built:	15 / 33 requirements		45%
Not built:	6 / 33 requirements		18%
Spec deviation:	1 / 33 requirements		3%

Launch-blocker gaps: 3 (age gate, zero-retention clarification, smart-home)

Post-launch gaps: 10 (CAP Orchestrator, audio cues, Vocspad, wallet, etc.)

Estimated engineering to reach 100% spec compliance:

- 2 developers × 8 weeks OR 1 developer × 16 weeks
- Estimated cost: \$24K–\$48K if contracted out

End of CUBIQO Complete Report (v2) — All sections consolidated.